High Performance Computing for Weather and Climate (HPC4WC)

Content: Caches and Data Locality
Lecturers: Oliver Fuhrer, Tobias Wicky
Block course 701-1270-00L
Summer 2020

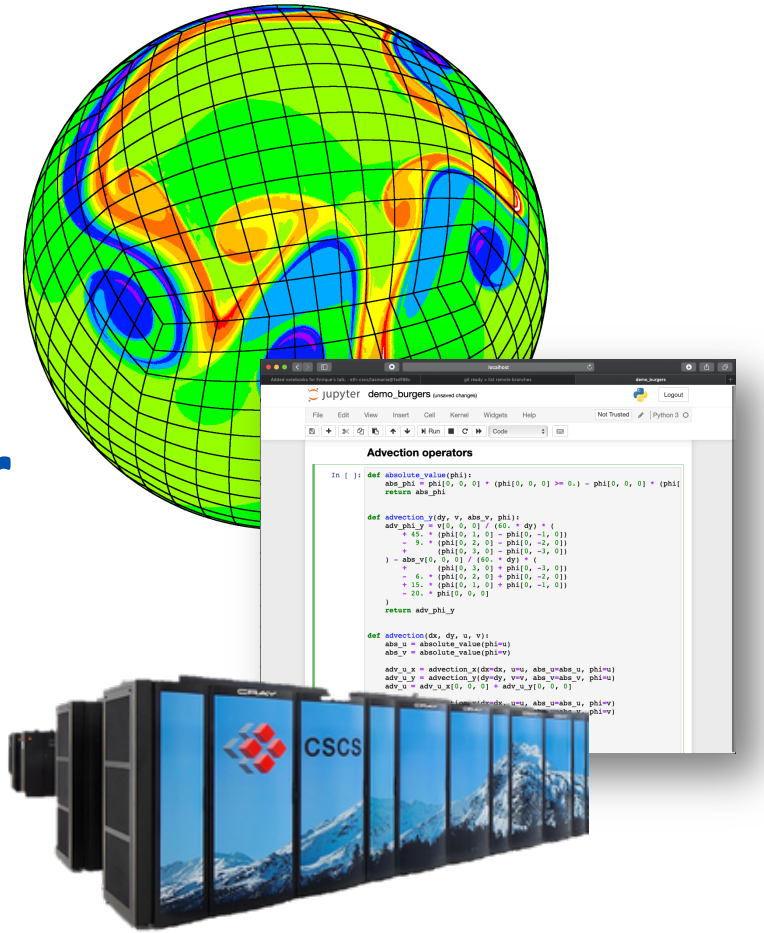

Von Neumann Architecture

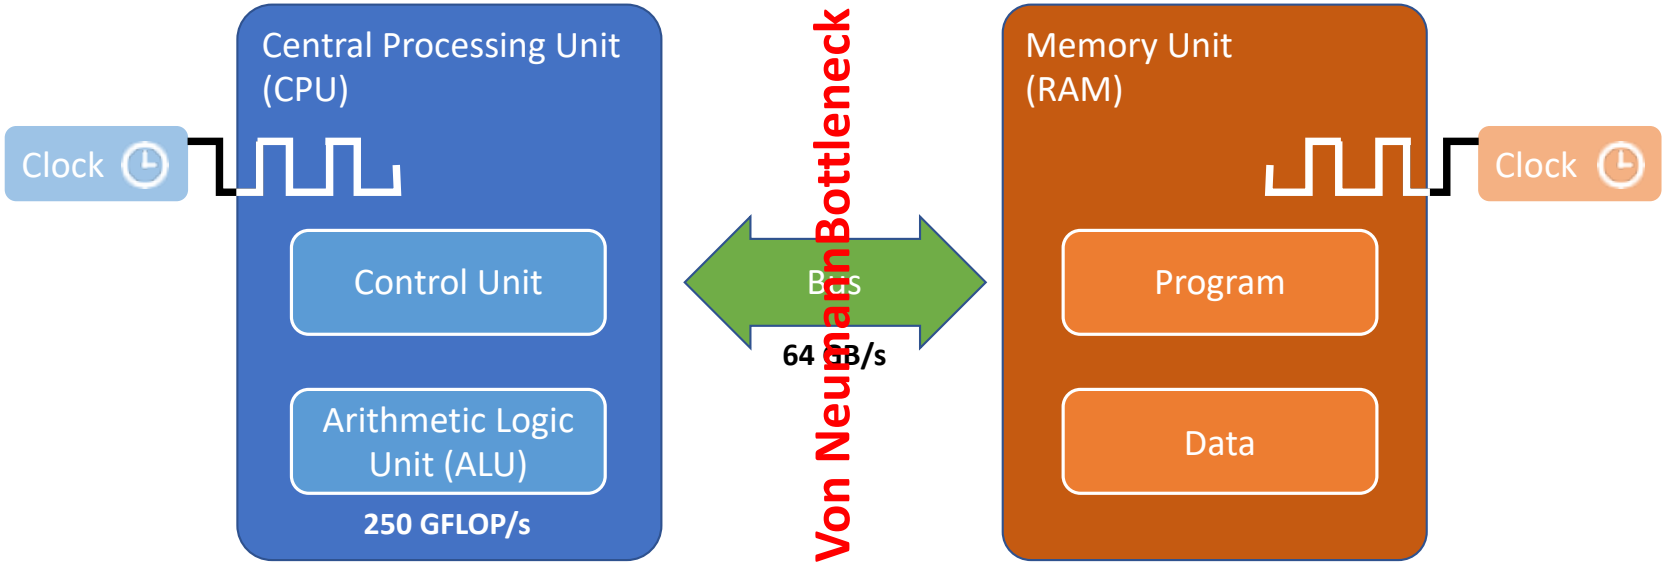

32 floating-point operations per load/store of a data value!

Stencil 2d

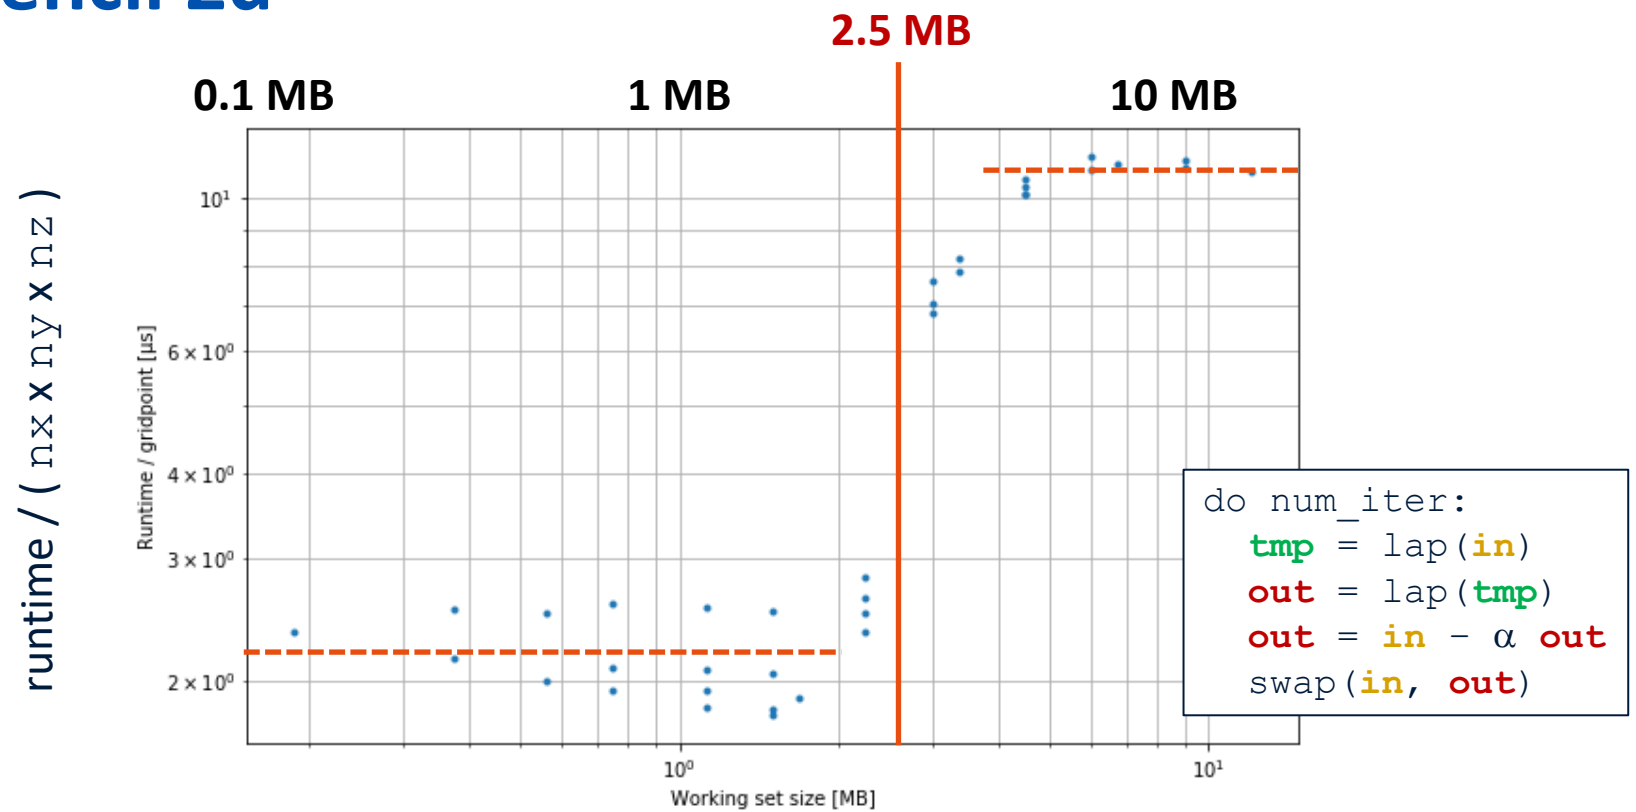

n = 3 fields x (nx x ny x nz) x 4 bytes

Learning goals

- Understand how data arrays are stored in computer memory
- Understand the implications of the cache hierarchy in a modern multi-core CPU
- Able to do basic data-locality optimizations (fusion, inlining) to improve performance

Cache Memory

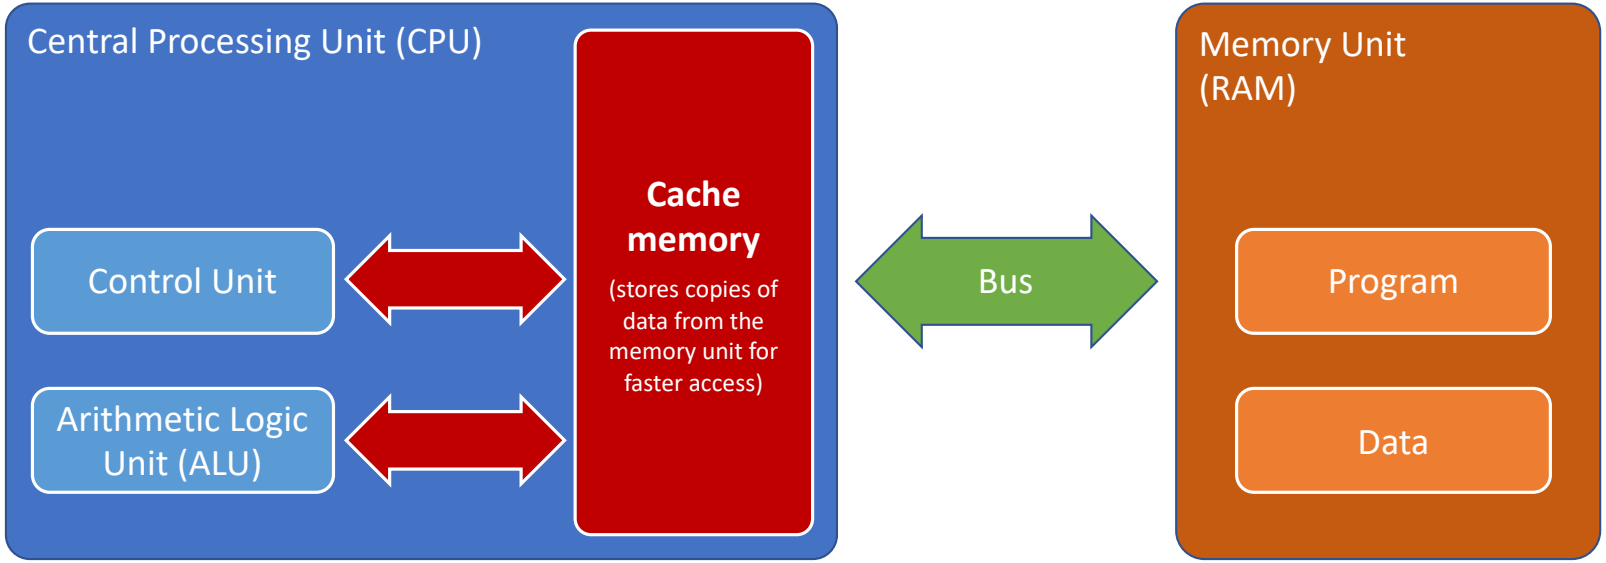

Cache is computer memory with short access time used for the storage of frequently or recently used data

Cache Mechanics

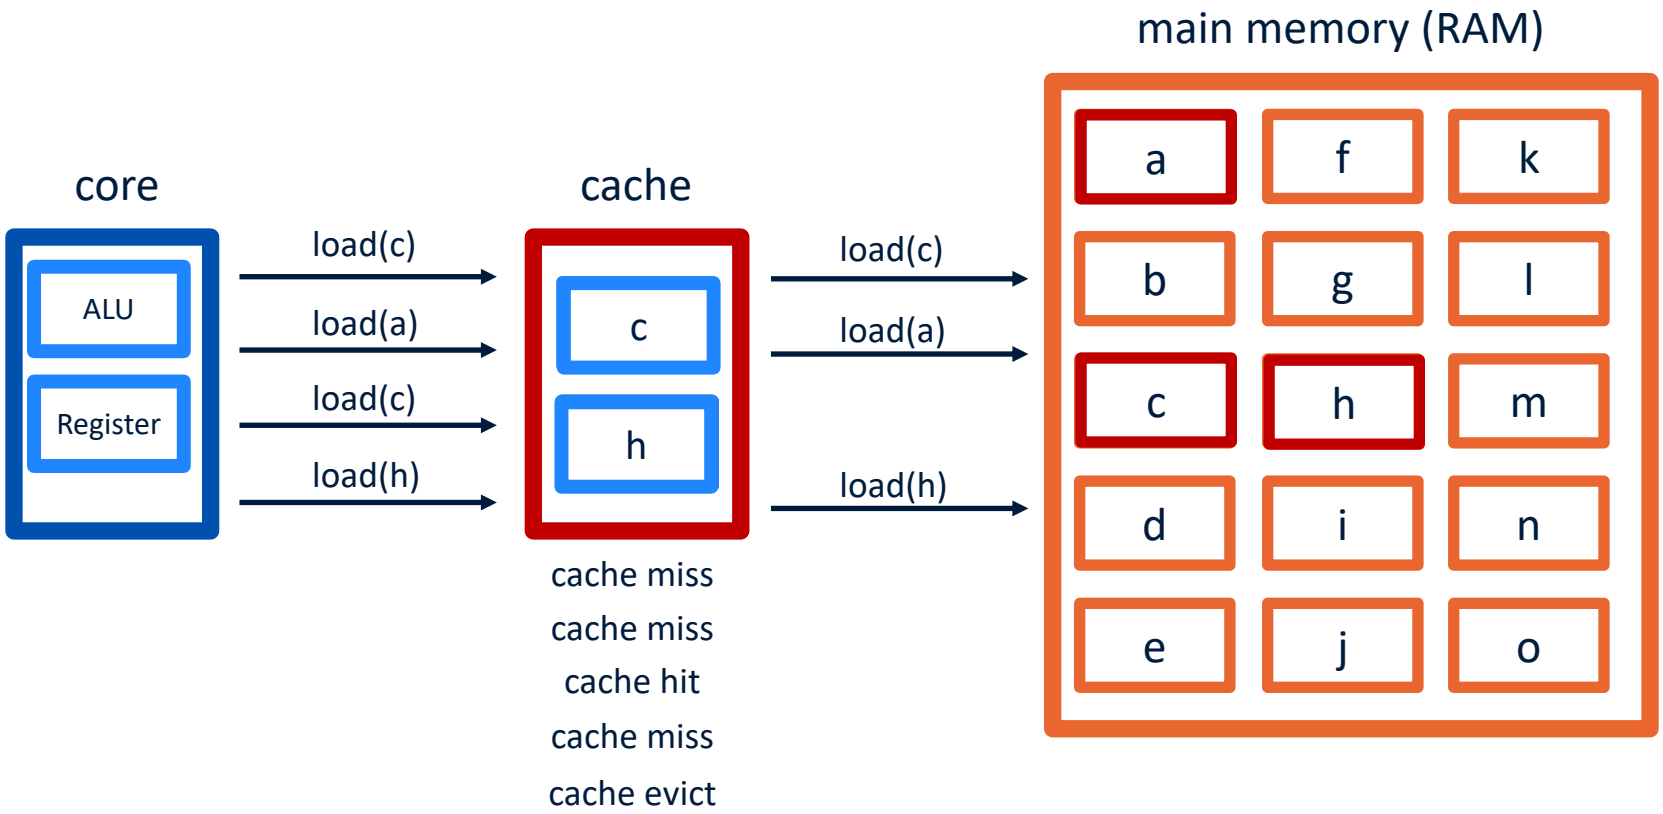

Cache Performance Metrics

Miss Rate

- fraction of memory references not found in cache (misses / accesses)
- miss rate = 1 – hit rate
- typical numbers (in percentages)
  - 3-10% for L1
  - can be quite small (< 1%) for L2, depending on size

Hit Time

- Time to deliver a line in the cache to the processor
  - includes time to determine whether the line is in the cache

Miss Penalty

- Additional time required because of a miss

How is data stored in memory?

```
real (kind=4) :: field(nx)
```

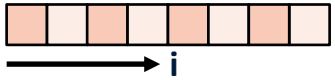

Stride in i-direction is 4 bytes

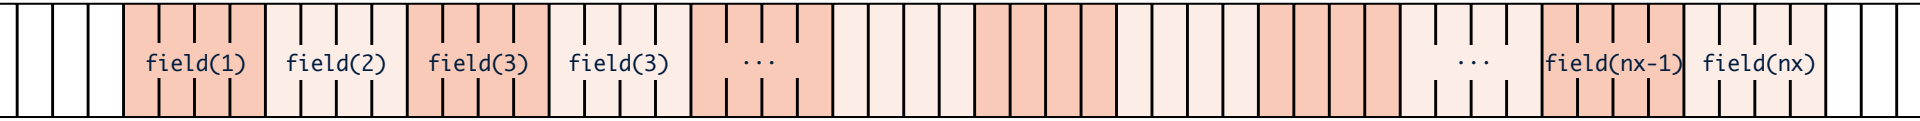


```
real (kind=4) :: field(nx, ny)
```

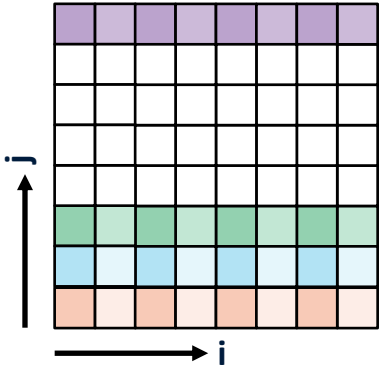

Stride in i-direction is 4 bytes
Stride in j-direction is 4 x `nx` bytes

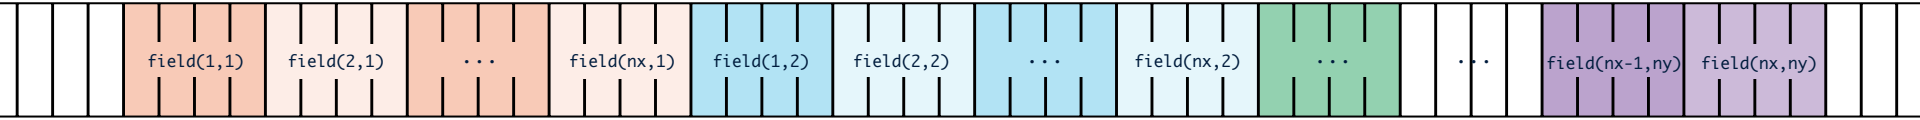

Why Caches Work

Locality: Programs tend to use data and instructions with addresses near or equal to those they have used recently

e.g. stencil computations

```
# weights literal constants
for j = 1, nj
  for i = 1, ni
    s_new(i, j) = &
      0.125 * s(i-1, j-1) + 0.25 * s(i, j-1) + 0.125 * s(i+1, j-1)  &
      0.25  * s(i-1, j  ) + 1.00 * s(i, j  ) + 0.25  * s(i+1, j  )  &
      0.125 * s(i-1, j+1) + 0.25 * s(i, j+1) + 0.125 * s(i+1, j+1)
```

Memory Hierarchy

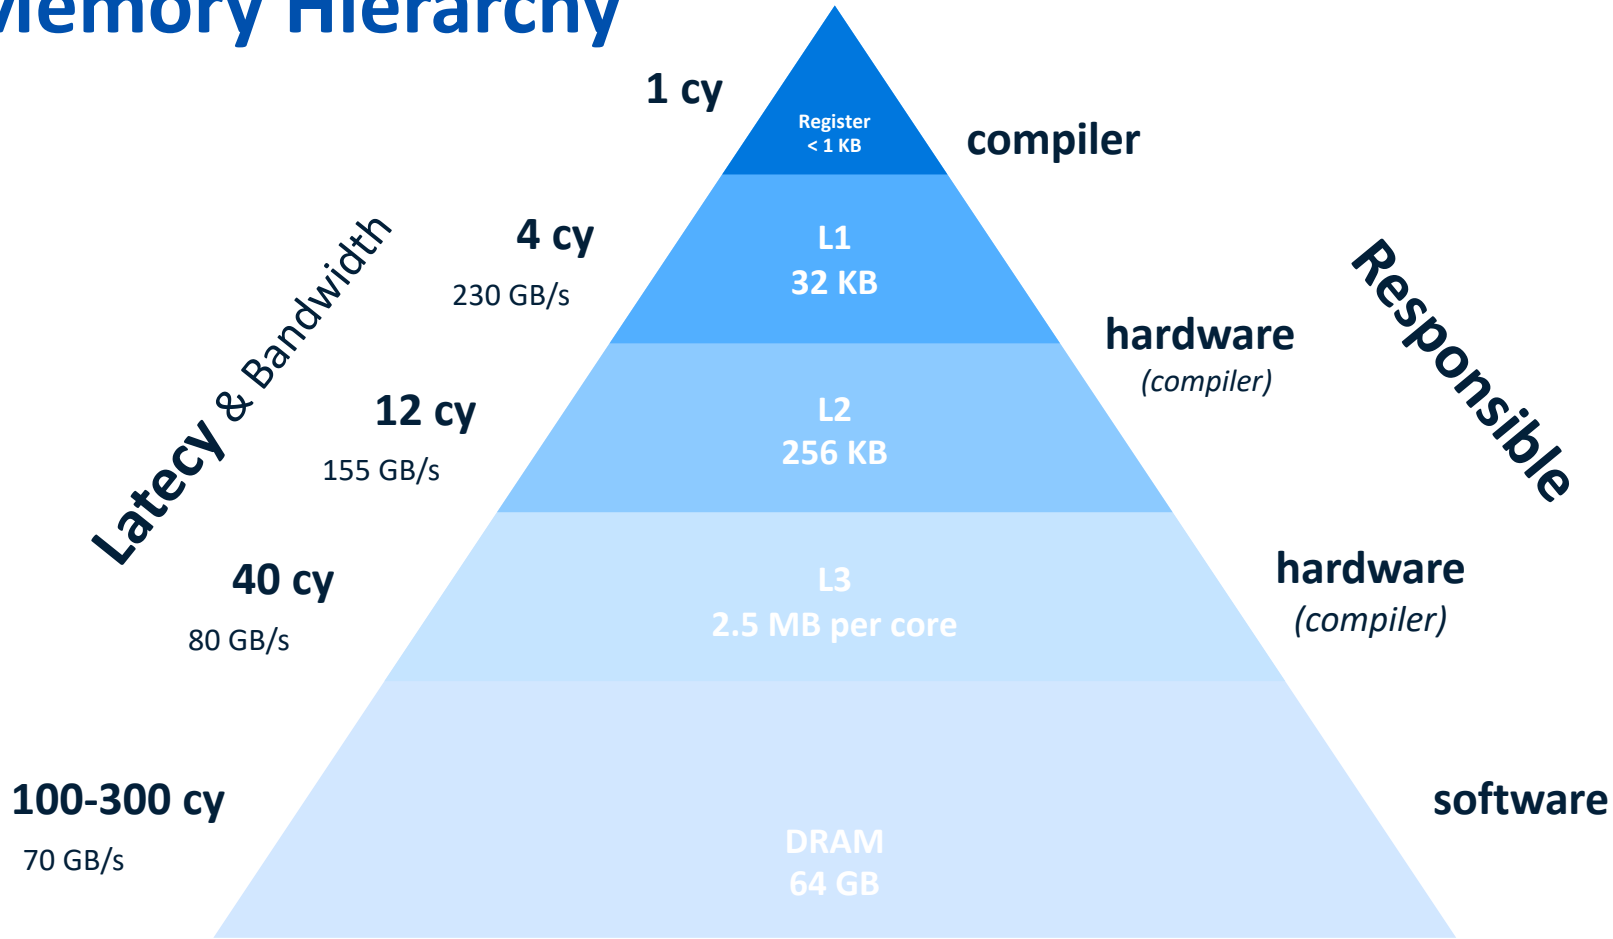

Cache Sizes (L1, L2, L3)

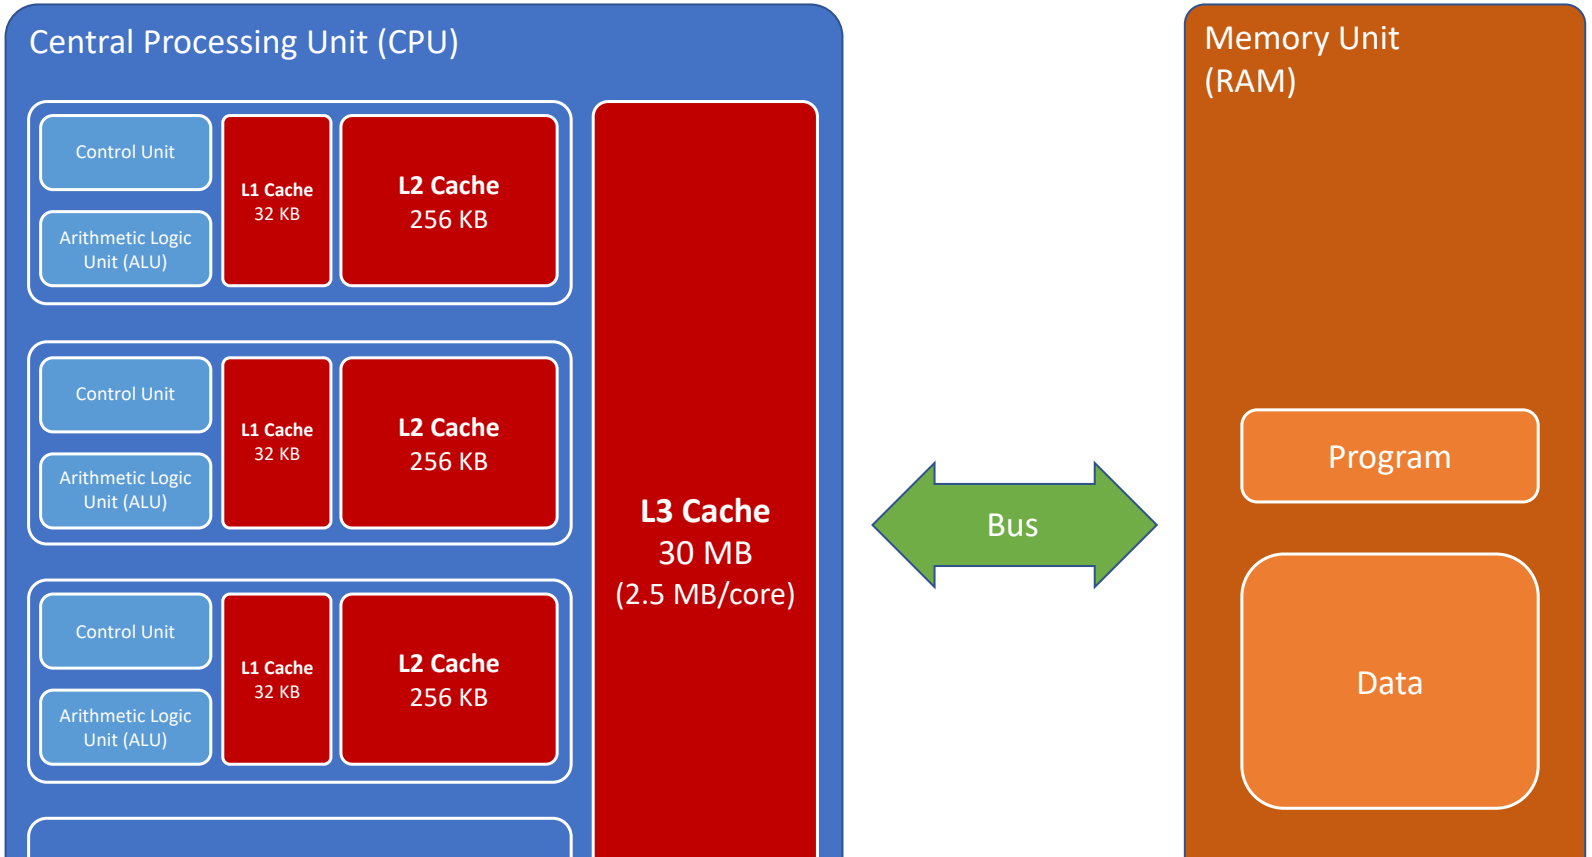

Stencil 2d

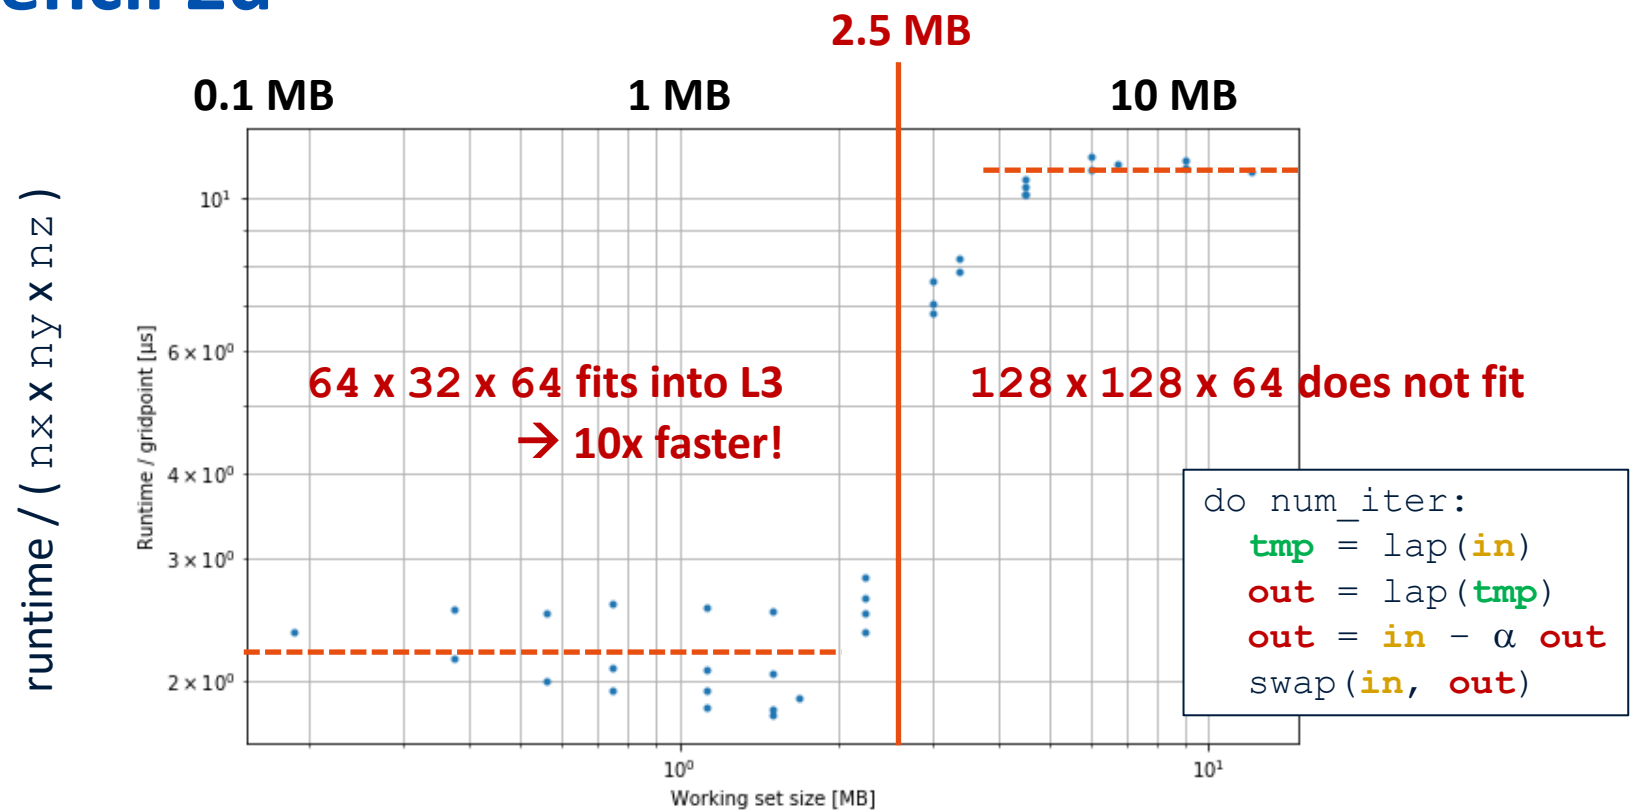

n = 3 fields x (nx x ny x nz) x 4 bytes

- Caches hold frequently requested data and are used to reduce memory access times.
- Modern CPUs have a hierarchy of caches (L1, L2, L3) of increasing size and access time.
- Data-locality optimizations aim to improve cache use (on all levels of the hierarchy) in order to improve performance.

Lab Exercises

01-roofline-model.ipynb

- Learn about performance metrics and how to compute theoretical peak values.
- Learn about arithmetic intensity and performance limiters.

02-stencil-program.ipynb

- Determine arithmetic intensity of a stencil program.
- Apply a performance profiling tool to gain insight into performance.
- Show limitations of the von Neumann model for understanding performance.

03-caches-data-locality.ipynb

- Learn about caches.
- Apply fusion in the stencil2d program and measure performance improvement.
- Apply inlining in the stencil2d program and measure performance improvement.

What is in (which) cache?

Reminder: L1 = 32 KB, L2 = 256 KB, L3 = 2.5 MB/core

```
real (kind=4) :: in(nx + 2*nh, ny + 2*nh, nz)
real (kind=4) :: tmp(nx + 2*nh, ny + 2*nh, nz)

do k = 1, nz
  do j = 1 + nh, ny + nh
    do i = 1 + nh, nx + nh
      tmp(i,j,k) = -4.0 * in(i,j,k) &
                   + in(i-1,j,k) + in(i+1,j,k) &
                   + in(i,j-1,k) + in(i,j+1,k)
```

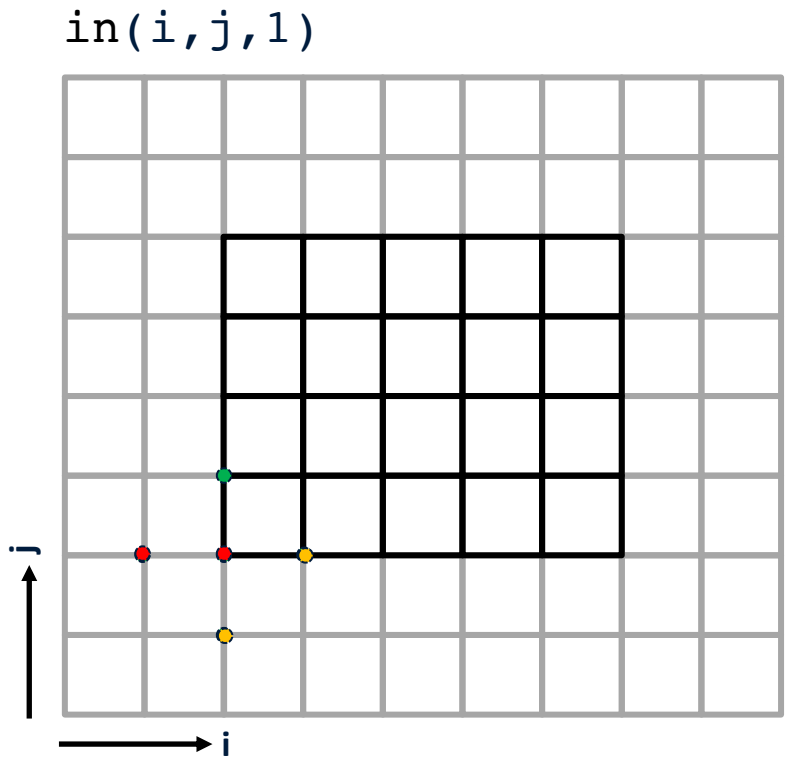

Stride in x-direction is 4 bytes

- Values `in(i,j,k)` and `in(i-1,j,k)` will always be in **L1 cache**

Stride in y-direction is approx. 4 x nx bytes

- If nx < 2048 we can retain 4 x-lines in **L1 cache**
- Then values at j and j-1 will be in cache from access to in(i,j+1,k)
- Only read `in(i,j+1,k)` and write `tmp(i,j,k)` from main memory!

Stride in z-direction is approx. 4 x nx x ny bytes

- For nx = ny = 128, the stride is 64 KB (too large for L1)

A full cube is approx. 4 x nx x ny x nz bytes (4 MB)

- For nx = ny = 128 and nz = 64 this is 4 MB
- If we start iterating again, `tmp` and `in` will be read from main memory!