

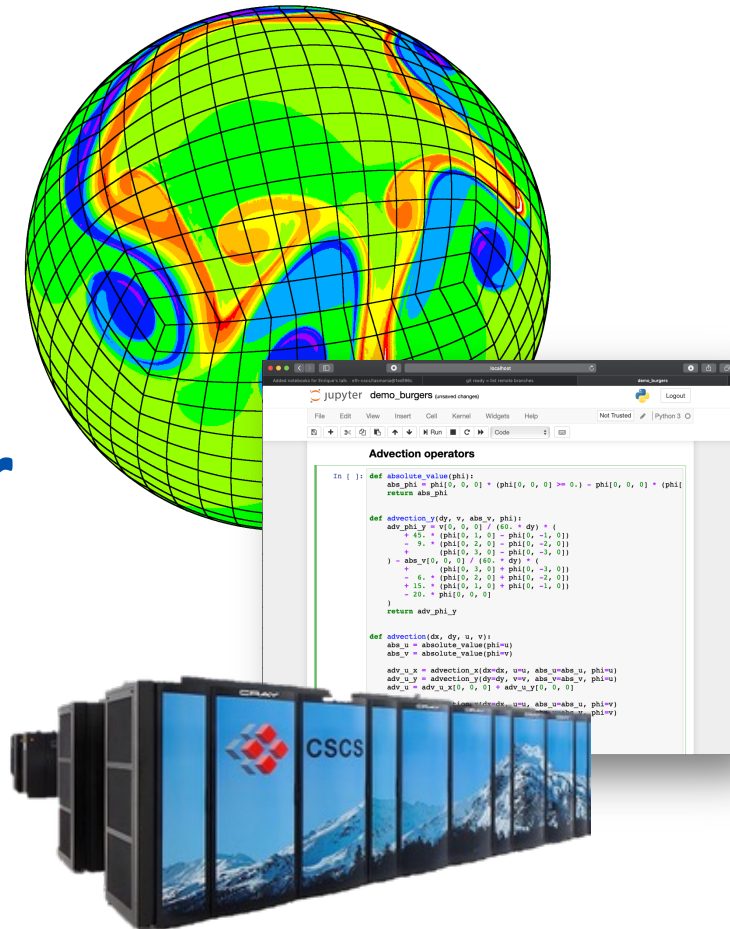
High Performance Computing for Weather and Climate (HPC4WC)

Content: Administrative

Lecturers: Oliver Fuhrer, Tobias Wicky, Stefano Ubbiali

Block course 701-1270-00L

Summer 2021



Nice to meet you!




Oliver Fuhrer
oliver.fuhrer@ethz.ch



Tobias Wicky
tobiasw@vulcan.com

Goals of course

- Understand high performance computing concepts relevant for weather and climate simulations
- Able to work with weather and climate simulation codes that run on large supercomputers



2016-08-11 18:00Z
258 Forecast Hours
FV3 3km

Approach

“ I Hear and I Forget, ”
I See and I Remember,
I Do and I Understand
(chinese proverb)

- Lectures that explain concepts and give context (*hear*).
- Demonstrations of the concepts being applied (*see*).
- Practical exercises and a work project (*do*).

Questions, please!

ASK QUESTIONS - BY JAKEPOSEY

WWW.TOONDOO.COM



Schedule

Monday	Motivation, stencil computations, memory hierarchy, lab environment
Tuesday	Shared memory parallelism, OpenMP, performance metrics
Wednesday	Distributed memory parallelism, domain-decomposition and halo-updates
Thursday	Hardware trends in supercomputing, GPU computing
Friday	High-level programming, domain-specific languages, wrapup

08:15 – 12:00	Morning session
12:00 – 13:30	<i>Lunch break</i>
13:30 – 17:00	Afternoon session

Questions?

Currently registered students

MSc Environmental Science (USYS)	8
MSc Computational Sciences and Engineering (MAVT)	6
PhD students (USYS, PHYS)	6
MSc Atmosphere and Climate Science (USYS)	4
MSc Physics (PHYS)	4
MSc Earth Sciences	2
MSc Mechanical Engineering (MAVT)	1
MSc Data Science	1
Total	32

Prerequisites

- **Fundamentals of numerical analysis and atmospheric modeling**
 - Basic partial differential calculus and finite difference methods.
 - e.g. ETH course “[Numerical methods in environmental physics](#)” or “[Numerical modeling of weather and climate](#)”
- **Experience in a programming language (C/C++, Fortran, Python, ...)**
 - We will read and write [Fortran](#), C++ and [Python](#) in this course.
- **Experience using command line interfaces in *nix environments (e.g., Unix, Linux)**
 - Familiar with work in the [command line shell](#) and the most commonly used shell commands.
 - Can logon to linux system via ssh and can work remotely on that system.
 - We will work on the [Piz Daint supercomputer](#) at the [Swiss National Supercomputing Center \(CSCS\)](#) in Lugano in this course.


If you think this course might not be suitable for you, contact us!

Questionnaire

(Work environment)

Have you worked with Jupyter Notebooks before?

23 responses




Questionnaire

(Programming)

How comfortable are you with this code example (CUDA)?

23 responses




Questionnaire

(Weather and Climate Modeling)

Do you have experience with a weather and/or climate model?

23 responses




Questionnaire

(High Performance Computing)

Have you seen a program deadlock due to bad communication patterns?

23 responses



Practicalities

- **All course material on GitHub repository** (slides, notebooks, codes, ...)
<https://github.com/ofuhrer/HPC4WC/>
- **Questions related to course in dedicated Slack workspace**
https://join.slack.com/t/hpc4wc-workspace/shared_invite/zt-1a0hd2f2s-X9oOK2DtpdcglroQPtfQAw
 - Generally, try to use public channels for questions since others probably have the same questions.
- **Lectures are not recorded**



How to earn credits (3 ECTS)

- **Attend the block course** (and participate actively!)
- **Work project**
 - Choose group and topic
 - Hand in working source code and report (max. 10 pages)
 - Projects will be graded
 - **Deadline: 31. August 2020**
- Credits are awarded if course attended and grade of work project ≥ 4.0
- Same rules apply for BSc, MSc, and PhD students


Work project

- **Work in groups of 2-3** (individual projects are strongly discouraged)
 - Programming is not a solitary art!
- **Topics will be presented on Friday**
 - If you prefer to choose your own, you are required to discuss with us beforehand
 - Each project must have a software development and performance evaluation part and has to be related to course material
- **Grading**
 - 25% correctness (compiles & runs, results correct, no bugs)
 - 25% quality (structure, clean code, comments, naming, tests, error handling)
 - 25% performance (depending on work project)
 - 25% report (maximum 10 pages)
- See [last year's projects](#) for examples


Example: Diffusion on a cubed-sphere grid



$t = 0$



$t = 100$



Questions?

Lab exercises

- Swiss National Supercomputing Centre <https://www.cscs.ch/>
- Piz Daint supercomputer (#23 on list of 500 largest supercomputers worldwide)



CSCS Accounts

- Everybody has a unique user name (classXXX) and password.
- **Do not share you login / pwd with anybody else.** Accounts with suspicious activities will be close down by CSCS immediately.
- **Change your password** immediately upon your first login to CSCS using the kpasswd command in a Terminal (see instructions).
- We have a shared quota of 3000 node hours for using the CSCS supercomputers for this block course.
 - Do not launch jobs with more than 1 node without checking with us first.
 - Do not leave your JupyterHub Server running if you don't need it.
- **Do not contact CSCS immediately** if you have trouble. Ask us or use the Slack workspace to get your issues resolved.
- Take a look at the [CSCS Code of Conduct](#)

JupyterHub

- Lab exercises will all be conducted on <https://jupyter.cscs.ch/>, the JupyterHub portal of CSCS.
- Interactive development and computing environment.
- If things get stuck or go wrong, it's always possible to “Stop Server” and “Launch Server” again.
- Jupyter notebooks auto-save and almost certainly no work will be lost.

The screenshot shows the JupyterHub interface on the CSCS website. The browser address bar shows `jupyter.cscs.ch`. The page header includes the CSCS logo and the ETH zürich logo. The user is logged in as `course01`. The interface has a navigation bar with `Home` and `Token` links. The main content area contains a form for launching a JupyterLab instance. The form includes fields for `Node Type` (set to `GPU`), `Nodes` (set to `1`), and `Duration (hr)` (set to `10`). There is an `Advanced options` button. Below this, there is a `Queue` dropdown (set to `Dedicated Queue (Max. 4 Nodes)`), a `Project Id` field, and a `JupyterLab Version` dropdown (set to `1.1.1`). The `Advanced Reservation` section has a text input field containing `climate`. The `Start IPyParallel Cluster with MPI Support?` section has a `No` radio button selected. The `Start Distributed Dask Cluster?` section has a `No` radio button selected. The `MPI Processes Per Node` field is set to `1`. The `Dask Tasks Per Node` field is set to `1`. At the bottom, there is a red button labeled `Launch JupyterLab`. Red annotations with numbers 1 through 5 are overlaid on the image: 1. Increase to 10 hours (pointing to the Duration field), 2. Click to expand (pointing to the Advanced options button), 3. Enter “climate” here (pointing to the Advanced Reservation text input), 4. Delete the “1” here (pointing to the MPI Processes Per Node field), and 5. Click to launch (pointing to the Launch JupyterLab button).

1. Increase to 10 hours


2. Click to expand

3. Enter “climate” here


4. Delete the “1” here

5. Click to launch


Literature & Links



Introduction to High Performance Computing for Scientists and Engineers, G. Hager and G. Wellein, CRC Press, 2011
([available online at ETH](#))



Parallel Computing, A. Grama, A. Gupta, G. Karypis, V. Kumar
([available free online](#))



Parallel Programming in MPI and OpenMP, V. Eijkhout
([Link to course](#))



Computer Organization and Design, D.H. Patterson and J.L. Hennessy ([available online at ETH](#))

Questions?

(10 minute break)