


High Performance Computing for Weather and Climate (HPC4WC)

Content: Administrative

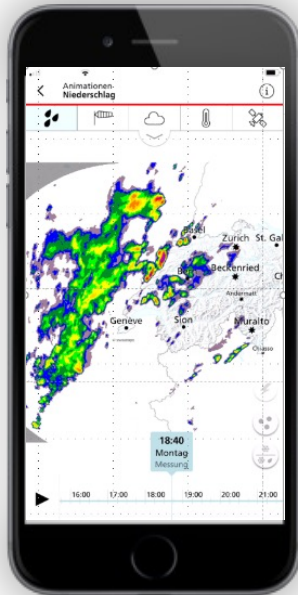
Lecturers: Oliver Fuhrer, Simon Adamov, Tobias Wicky

Block course 701-1270-00L

Summer 2024



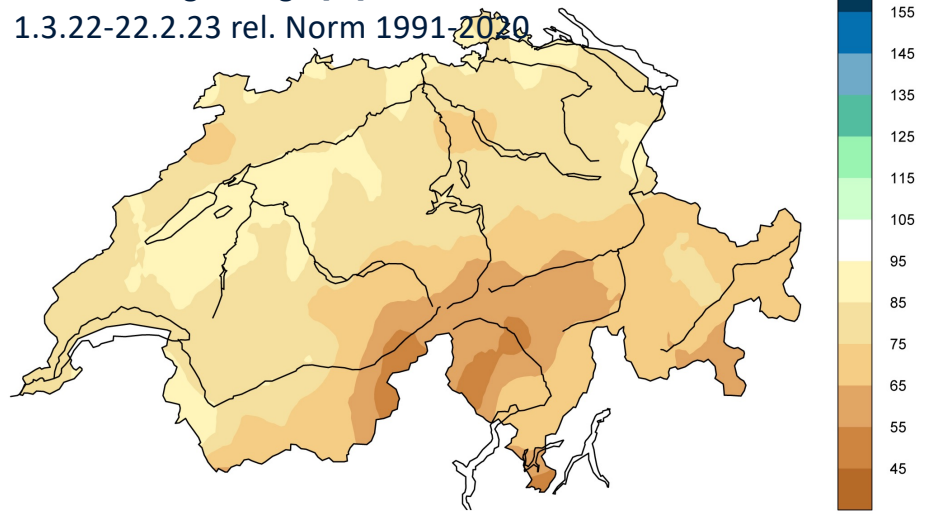
July 13, 2021 2am in the morning...




What's behind the warnings for such extreme events?

Niederschlagsmenge [%]

1.3.22-22.2.23 rel. Norm 1991-2020



How do we know if this will be the «norm» in the future?



2016-08-11 18:00Z
258 Forecast Hours
FV3 3km

Weather and climate models run on supercomputers



Goals of course

- Understand high performance computing concepts relevant for weather and climate simulations
- Able to work with weather and climate simulation codes that run on large supercomputers

Approach

“ I *hear*, and I forget. ”
I *see*, and I remember.
I *do*, and I understand.
(chinese proverb)

- Lectures that explain concepts and give context (*hear*).
- Demonstrations of the concepts being applied (*see*).
- Practical exercises and a work project (*do*).

Questions, please!

ASK QUESTIONS - BY JAKEPOSEY



Schedule

Monday	Motivation, stencil computations, memory hierarchy, lab environment
Tuesday	Shared memory parallelism, OpenMP, performance metrics
Wednesday	Distributed memory parallelism, domain-decomposition and halo-updates
Thursday	Hardware trends in supercomputing, GPU computing
Friday	High-level programming, domain-specific languages, wrapup

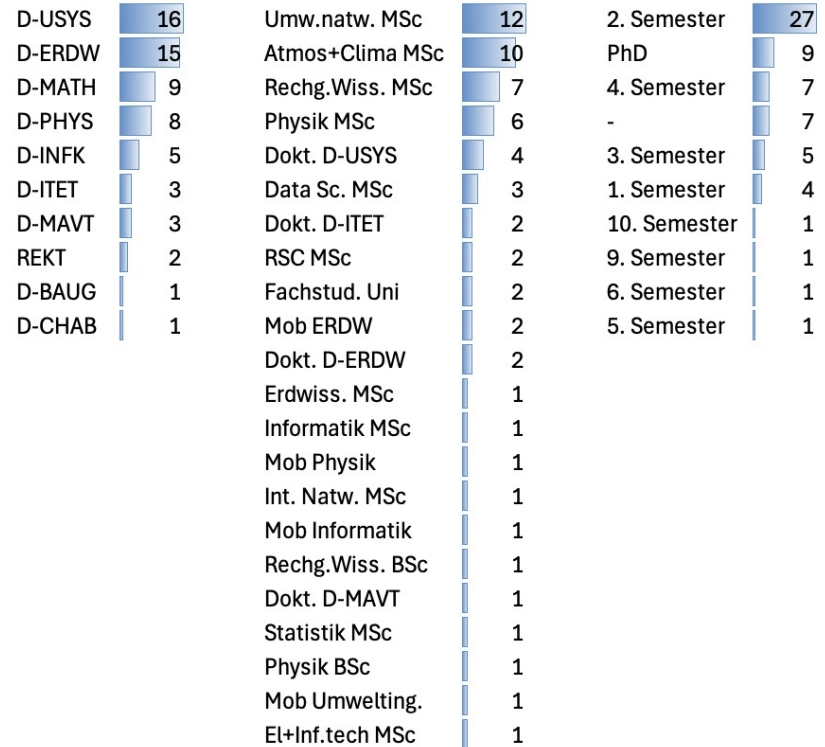
08:15 – 12:00	Morning session
12:00 – 13:30	<i>Lunch break</i>
13:30 – 17:00	Afternoon session

Currently registered students (eDoz)

We are currently
53 registered participants.

Too much for the format
of the course.

We only have 48 CSCS
accounts.



Prerequisites

- **Fundamentals of numerical analysis and atmospheric modeling**
 - Basic partial differential calculus and finite difference methods.
 - e.g. ETH course “[Numerical methods in environmental physics](#)” or “[Weather and climate models](#)”
- **Experience in a programming language (C/C++, Fortran, Python, ...)**
 - We will read and write [Fortran](#), C++ and [Python](#) in this course.
- **Experience using command line interfaces in *nix environments (e.g., Unix, Linux)**
 - Familiar with work in the [command line shell](#) and the most commonly used shell commands.
 - Can logon to linux system via ssh and can work remotely on that system.
 - We will work on the [Piz Daint supercomputer](#) at the [Swiss National Supercomputing Center \(CSCS\)](#) in Lugano in this course.

If you think this course might not be suitable for you, talk to us!

Quick Poll



Linux / Unix, Terminal



Programming (Python, Fortran, C/C++, ...), Compilation, Debugger



Jupyter Notebooks



PDE, Stencil, Weather Model



Numpy, OpenMP, MPI, CUDA

Practicalities

- **All course material on GitHub repository** (slides, notebooks, codes, ...)
<https://github.com/ofuhrer/HPC4WC/>
- **Questions related to course and projects in dedicated Slack workspace**
https://join.slack.com/t/hpc4wc-workspace/shared_invite/zt-1a0hd2f2s-X9oOK2DtpdcgIroQPtfQAw
- Generally, try to use public channels for questions since others probably have the same questions.
- **Lectures are not recorded**



How to earn credits (3 ECTS)

- **Attend the block course** (and participate actively!)
- **Work project**
 - Choose group and topic
 - Hand in working source code and report (max. 10 pages)
 - Projects will be graded
 - **Deadline: 31. August 2024**
- Credits are awarded if course attended and grade of work project ≥ 4.0
- Same rules apply for BSc, MSc, and PhD students


Work project

- **Work in groups of 3-4** (no individual projects)
 - Programming is not a solitary art!
- **Topics will be presented on Wednesday**
 - Have an idea of your own? Cool, discuss with us to shape it into a project!
 - Each project must have a software development and performance evaluation part and has to be related to course material
- **Grading**
 - 25% correctness (compiles & runs, results correct, no bugs)
 - 25% quality (structure, clean code, comments, naming, tests, error handling)
 - 25% performance (depending on work project)
 - 25% report (maximum 10 pages)
- See [last year's projects](#) for examples

Examples

Shallow water equations on a toroidal planet
in the domain-specific language GT4Py

Killian P. Brennan, Dana Grund, Joren Janzing, and Franco Lee
Course Project: High-Performance Computing for Weather and Climate
August 2023



CACHE HIERARCHY AND LOOP OPTIMIZATIONS

Report

Basil Ruch*
ETH Zurich, Switzerland
Department of Mathematics

David Strassmann†
ETH Zurich, Switzerland
Department of Mathematics

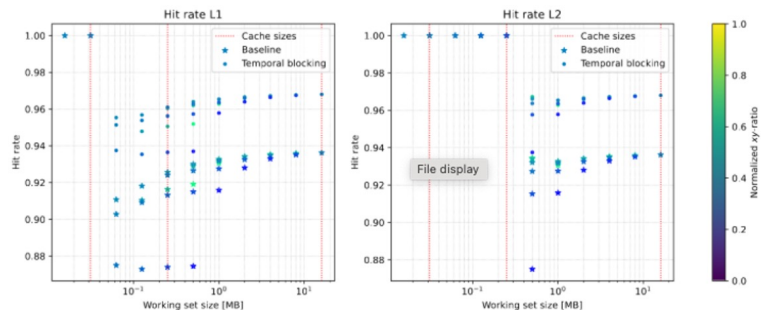


Figure 4: Recorded hit rates in the L1 (left) and L2 (right) cache for the 2D stencil. Star markers represent the baseline and circles the temporal blocking algorithm. The color indicates the equality of the two dimensions: a value of 0.5 means that x and y dimension are equal.

Lab exercises

- Swiss National Supercomputing Centre <https://www.cscs.ch/>
- Piz Daint supercomputer (#43 on [list of 500 largest supercomputers](#) worldwide)



CSCS Accounts

- Everybody has a unique user name (classXXX) and password.
- **Do not share you login / pwd with anybody else.** Accounts with suspicious activities will be closed down by CSCS immediately.
- **Change your password** immediately upon your first login to CSCS using the `passwd` command in a Terminal (see instructions).
- We have a shared quota of 4000 node hours for using the CSCS supercomputers for this block course.
 - Do not launch jobs with more than 1 node without checking with us first.
 - Do not leave your JupyterHub Server running if you don't need it.
- **Do not contact CSCS first** if you have trouble. Ask us or use the Slack workspace to get your issues resolved.
- Take a look at the [CSCS Code of Conduct](#)

JupyterHub

- Lab exercises will all be conducted on <https://jupyter.cscs.ch/>, the JupyterHub portal of CSCS.
- Interactive development and computing environment.
- If things get stuck or go wrong, it's always possible to “Stop Server” and “Launch Server” again.
- Jupyter notebooks auto-save and *almost* certainly no work will be lost.


The screenshot shows the JupyterLab launch interface with several elements circled in red and numbered 1 through 4:

- 1. Increase duration:** A red circle around the number '4' in the 'Duration (hr)' field.
- 2. Click:** A red circle around the 'Advanced options' toggle button.
- 3. Enter reservation:** A red circle around the text 'HPC4WC' in the 'Advanced Reservation' field.
- 4. Click:** A red circle around the 'Launch JupyterLab' button at the bottom.


The interface includes the following fields and controls:

- Node Type:** A dropdown menu currently set to 'GPU'.
- Nodes:** A numeric input field set to '1' with minus and plus buttons.
- Duration (hr):** A numeric input field set to '4' with minus and plus buttons.
- Queue:** A dropdown menu currently set to 'Dedicated Queue (Max. 5 Nodes)'.
- Project Id:** A text input field with the placeholder '(leave empty for default)'.
- Advanced Reservation:** A text input field containing 'HPC4WC'.
- JupyterLab Version:** A dropdown menu currently set to '3.2'.
- Launch JupyterLab:** A large red button at the bottom.


Literature & Links



Introduction to High Performance Computing for Scientists and Engineers, G. Hager and G. Wellein, CRC Press, 2011
(available online at [ETH](#))



Parallel Computing, A. Grama, A. Gupta, G. Karypis, V. Kumar
(available free online)



Parallel Programming in MPI and OpenMP, V. Eijkhout
([Link to course](#))



Computer Organization and Design, D.H. Patterson and J.L. Hennessy (available online at [ETH](#))