# High Performance Computing for Weather and Climate (HPC4WC)

Content: Caches and Data Locality
Lecturers: Oliver Fuhrer
Block course 701-1270-00L
Summer 2020

Learning goals

- Understand how data arrays are stored in computer memory

- Understand the implications of the cache hierarchy in a modern multi-core CPU

- Able to do basic data-locality optimizations (fusion, inlining) to improve performance

Cache Memory

Cache is computer memory with short access time used for the storage of frequently or recently used data

Cache Mechanics

core

main memory (RAM)

cache

a	f	k
b	g	l
c	h	m
d	i	n
e	j	o

ALU

Register

load(c)

load(a)

load(c)

load(h)

c

h

load(c)

load(a)

load(h)

cache miss

cache miss

cache hit

cache miss

cache evict

Cache Performance Metrics

Miss Rate

- fraction of memory references not found in cache (misses / accesses)
- miss rate = 1 – hit rate
- typical numbers (in percentages)
  - 3-10% for L1
  - can be quite small (< 1%) for L2, depending on size

Hit Time

- Time to deliver a line in the cache to the processor
  - includes time to determine whether the line is in the cache

Miss Penalty

- Additional time required because of a miss

How is data stored in memory?

`real (kind=4) :: field(nx)`

Stride in i-direction is 4 bytes

| field(1) | field(2) | field(3) | field(3) | ... | | | ... | field(nx-1) | field(nx) |

`real (kind=4) :: field(nx, ny)`

Stride in i-direction is 4 bytes
Stride in j-direction is 4 x `nx` bytes

| field(1,1) | field(2,1) | ... | field(nx,1) | field(1,2) | field(2,2) | ... | field(nx,2) | ... | ... | field(nx-1,ny) | field(nx,ny) |

Why Caches Work

Locality: Programs tend to use data and instructions with addresses near or equal to those they have used recently

e.g. stencil computations

```
# weights literal constants
for j = 1, nj
  for i = 1, ni
    s_new(i, j) = &
      0.125 * s(i-1, j-1) + 0.25 * s(i, j-1) + 0.125 * s(i+1, j-1)  &
      0.25  * s(i-1, j  ) + 1.00 * s(i, j  ) + 0.25  * s(i+1, j  )  &
      0.125 * s(i-1, j+1) + 0.25 * s(i, j+1) + 0.125 * s(i+1, j+1)
```

Memory Hierarchy

Latecy & Bandwidth

1 cy

4 cy
230 GB/s

12 cy
155 GB/s

40 cy
80 GB/s

100-300 cy
70 GB/s

Register
< 1 KB

**L1
32 KB**

**L2
256 KB**

**L3
2.5 MB per core**

**DRAM
64 GB**

compiler

hardware
(compiler)

hardware
(compiler)

software

Responsible

Cache Sizes (L1, L2, L3)

Central Processing Unit (CPU)

Control Unit

Arithmetic Logic Unit (ALU)

L1 Cache 32 KB

L2 Cache 256 KB

Control Unit

Arithmetic Logic Unit (ALU)

L1 Cache 32 KB

L2 Cache 256 KB

Control Unit

Arithmetic Logic Unit (ALU)

L1 Cache 32 KB

L2 Cache 256 KB

L3 Cache 30 MB (2.5 MB/core)

Bus

Memory Unit (RAM)

Program

Data

Silicon Real Estate

A large fraction of die real estate is dedicated to cache memory

Stencil 2d

2.5 MB

0.1 MB **1 MB** **10 MB**

64 x 32 x 64 fits into L3
→ 10x faster!

128 x 128 x 64 does not fit

runtime / (nx x ny x nz)

Runtime / gridpoint [μs]

10^1

6×10^0

4×10^0

3×10^0

2×10^0

10^0 10^1

Working set size [MB]

```
do num_iter:
    tmp = lap(in)
    out = lap(tmp)
    out = in - α out
    swap(in, out)
```
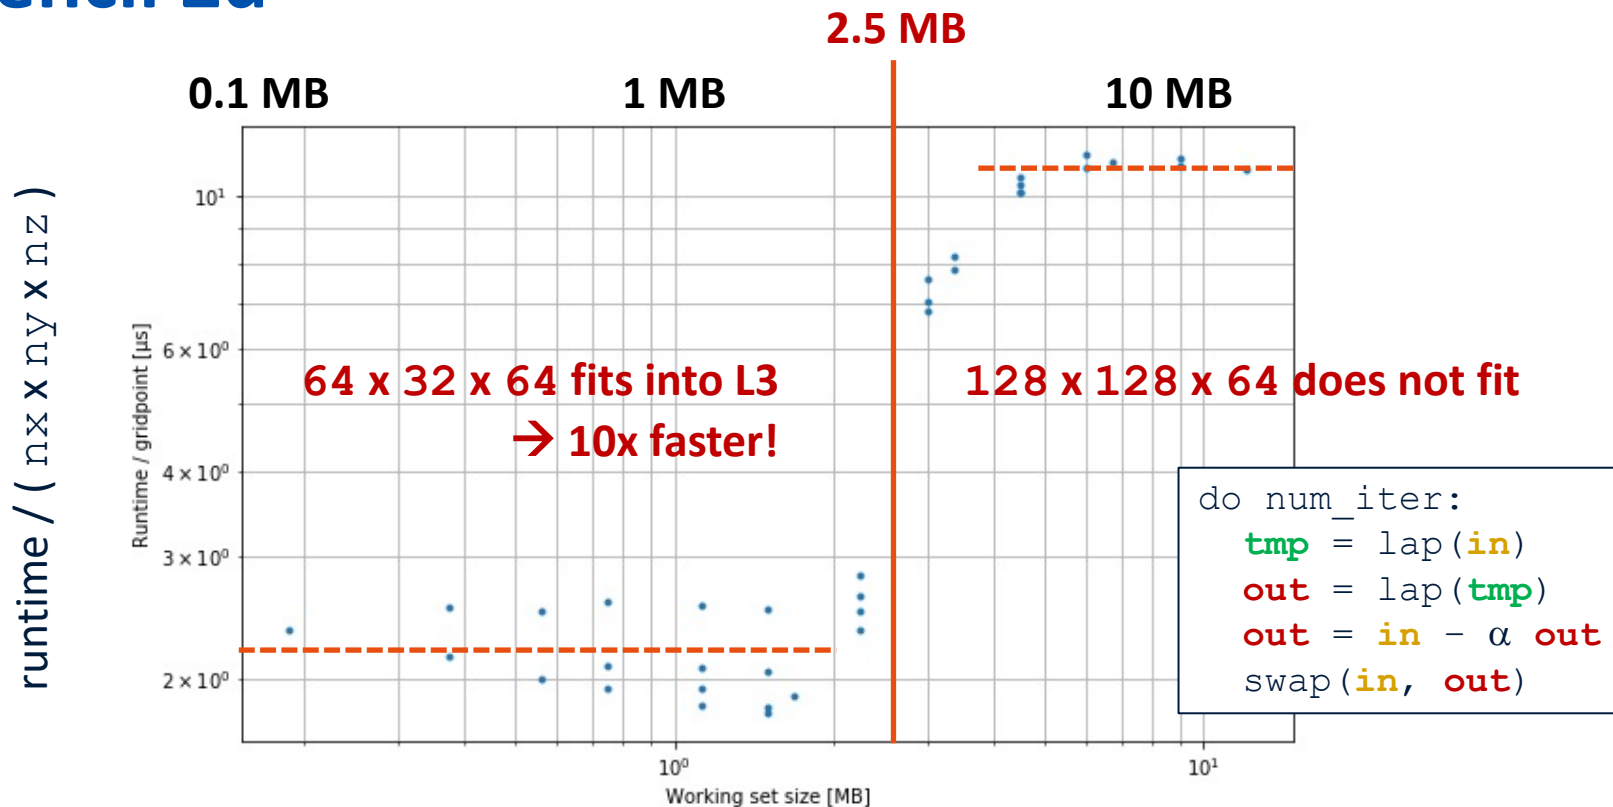
n = 3 fields x (nx x ny x nz) x 4 bytes

What is in (which) cache?

`in(:,:,1)`

```
real (kind=4) :: in(nx + 2*nh, ny + 2*nh, nz)
real (kind=4) :: tmp(nx + 2*nh, ny + 2*nh, nz)

do k = 1, nz
  do j = 1 + nh, ny + nh
    do i = 1 + nh, nx + nh
      tmp(i,j,k) = -4.0 * in(i,j,k) &
                   + in(i-1,j,k) + in(i+1,j,k) &
                   + in(i,j-1,k) + in(i,j+1,k)
```

Stride in x-direction is 4 bytes
- Values `in(i,j,k)` and `in(i-1,j,k)` are probably in **L1 cache**

Stride in y-direction is approx. 4 x nx bytes
- For nx=128, the j-stride is ~ 512 B
- If nx < 2048 we can retain approx. 4 i-lines in **L1 cache** (32 KB)
- For nx=128 the loads of `in(i+1,j,k)` and `in(i,j-1,k)` will be in L1
- Only read `in(i,j+1,k)` and write `tmp(i,j,k)` from main memory!

Stride in z-direction is approx. 4 x nx x ny bytes
- For nx=ny=128, the k-stride is ~64 KB
- **If we do multiple iterations over the ij-plane (k-blocking) we are in L2!**

A full cube is approx. 4 x nx x ny x nz bytes (4 MB)
- For nx=ny=128 and nz=64 this is ~4 MB
- If we do multiple iterations over ijk, `tmp` and `in` will be read from main memory!

Summary

- Caches hold frequently requested data and are used to reduce memory access times.

- Modern CPUs have a hierarchy of caches (L1, L2, L3) of increasing size and access time.

- Data-locality optimizations aim to improve cache use (on all levels of the hierarchy) in order to improve performance.

Lab Exercises

01-roofline-model.ipynb
- Learn about performance metrics and how to compute theoretical peak values.
- Learn about arithmetic intensity and performance limiters.

02-stencil-program.ipynb
- Determine arithmetic intensity of a stencil program.
- Apply a performance profiling tool to gain insight into performance.
- Show limitations of the von Neumann model for understanding performance.

03-caches-data-locality.ipynb
- Learn about caches.
- Apply fusion in the stencil2d program and measure performance improvement.
- Apply inlining in the stencil2d program and measure performance improvement.