

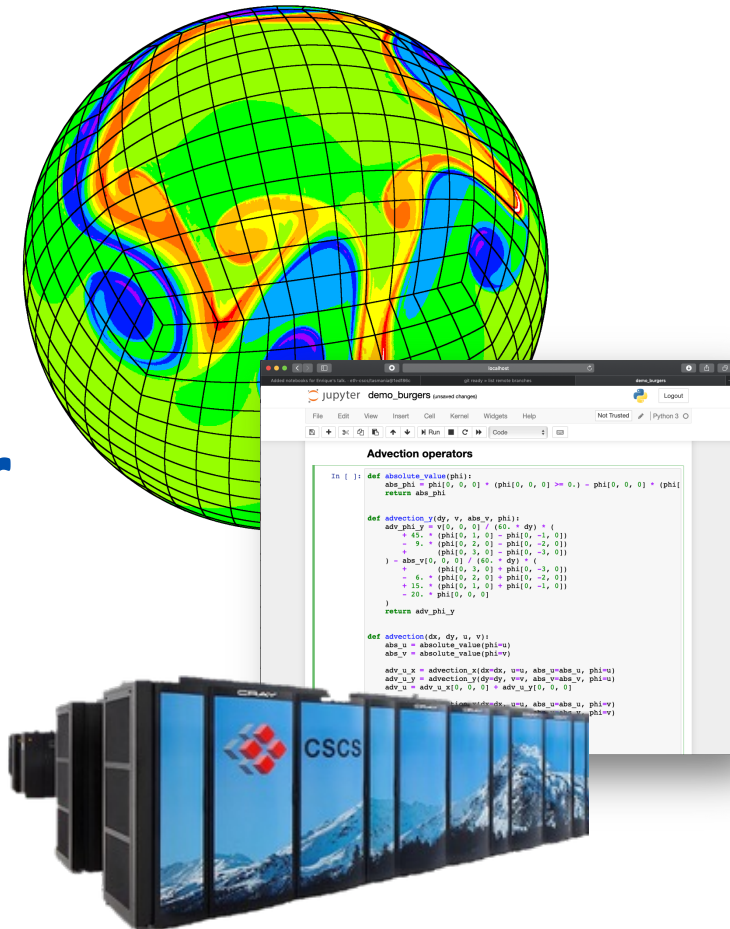
High Performance Computing for Weather and Climate (HPC4WC)

Content: Graphics Processing Units

Lecturer: Oliver Fuhrer

Block course 701-1270-00L

Summer 2024




Learning goals

- Understand why specialized hardware such as GPUs is become the new “normal”
- Learn how to program a GPU using a high-level programming language
- Understand potential and difficulties of GPU-computing

How does the landscape of HPC look today?


Moore's Law (1965)

- "The number of transistors in a dense integrated circuit will double every two years"




Original data up to the year 2010 collected and plotted by M. Horowitz, F. Labonte, O. Shacham, K. Olukotun, L. Hammond, and C. Batten
New plot and data collected for 2010-2021 by K. Rupp

The End of General Purpose Computing




Distance between Si-atoms is 0.5 nm!

Chip Design Costs



How does performance of our machines behave




So why are we still ok?

Rank	System	Cores	Rmax (PFlop/s)	Rpeak (PFlop/s)	Power (kW)
1	Frontier - HPE Cray EX235a, AMD Optimized 3rd Generation EPYC 64C 2GHz, AMD Instinct MI250X, Slingshot-11, HPE DOE/SC/Oak Ridge National Laboratory United States	8,699,904	1,194.00	1,679.82	22,703
AMD GPU					
2	Supercomputer Fugaku - Supercomputer Fugaku, A64FX 48C 2.2GHz, Tofu interconnect D, Fujitsu RIKEN Center for Computational Science Japan	7,630,848	442.01	537.21	29,899
3	LUMI - HPE Cray EX235a, AMD Optimized 3rd Generation EPYC 64C 2GHz, AMD Instinct MI250X, Slingshot-11, HPE EuroHPC/CSC Finland	2,220,288	309.10	428.70	6,016
AMD GPU					
4	Leonardo - BullSequana XH2000, Xeon Platinum 8358 32C 2.6GHz, NVIDIA A100 SXM4 64 GB, Quad-rail NVIDIA HDR100 Infiniband, Atos EuroHPC/CINECA Italy	1,824,768	238.70	304.47	7,404
NVIDIA GPU					
5	Summit - IBM Power System AC922, IBM POWER9 22C 3.07GHz, NVIDIA Volta GV100, Dual-rail Mellanox EDR Infiniband, IBM DOE/SC/Oak Ridge National Laboratory United States	2,414,592	148.60	200.79	10,096
NVIDIA GPU					

Rank	System	Cores	Rmax (PFlop/s)	Rpeak (PFlop/s)	Power (kW)
6	Sierra - IBM Power System AC922, IBM POWER9 22C 3.1GHz, NVIDIA Volta GV100, Dual-rail Mellanox EDR Infiniband, IBM / NVIDIA / Mellanox DOE/NNSA/LLNL United States	1,572,480	94.64	125.71	7,438
NVIDIA GPU					
7	Sunway TaihuLight - Sunway MPP, Sunway SW26010 260C 1.45GHz, Sunway, NRCPC National Supercomputing Center in Wuxi China	10,649,600	93.01	125.44	15,371
8	Perlmutter - HPE Cray EX235n, AMD EPYC 7763 64C 2.45GHz, NVIDIA A100 SXM4 40 GB, Slingshot-10, HPE DOE/SC/LBNL/NERSC United States	761,856	70.87	93.75	2,589
NVIDIA GPU					
9	Selene - NVIDIA DGX A100, AMD EPYC 7742 64C 2.25GHz, NVIDIA A100, Mellanox HDR Infiniband, Nvidia NVIDIA Corporation United States	555,520	63.46	79.22	2,646
NVIDIA GPU					
10	Tianhe-2A - TH-IVB-FEP Cluster, Intel Xeon E5-2692v2 12C 2.2GHz, TH Express-2, Matrix-2000, NUDT National Super Computer Center in Guangzhou China	4,981,760	61.44	100.68	18,482
Xeon Phi					


Specialized Chips are on the Rise!



Google's TPU
(e.g. machine learning)




FPGA
(e.g. bitcoin mining)



GPU
(e.g. gaming)

Dennard Scaling

- "If the transistor density doubles, power consumption (with twice the number of transistors) stays the same."




Performance / Watt

Intel Xeon E5-2690 v3 + DRAM



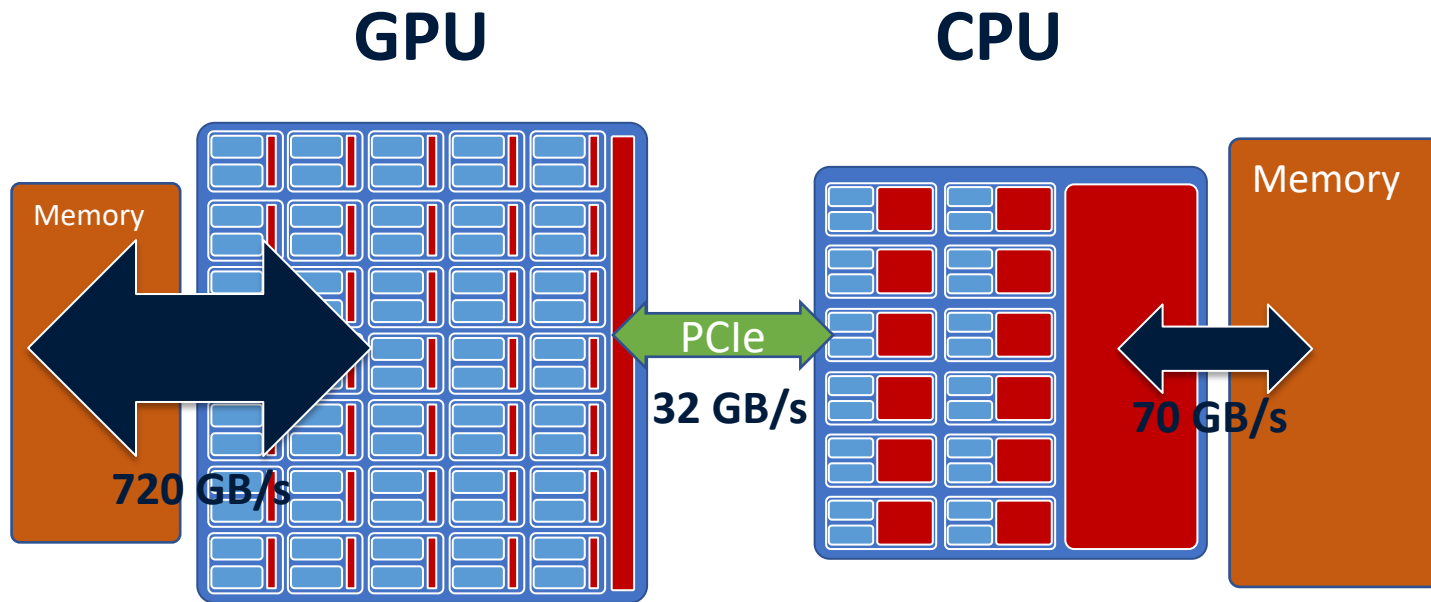
~ 200 W 0.5 TFLOP/s 70 GB/s

NVIDIA Tesla P100




~ 300 W 5.3 TFLOP/s 720 GB/s

Node Architecture




Crucial to minimize memory transfers between CPU and GPU!

Hybrid Supercomputer




Performance / Watt

Intel Xeon E5-2690 v3 + DRAM




~ 200 W 0.5 TFLOP/s 70 GB/s

NVIDIA Tesla P100



~ 300 W 5.3 TFLOP/s 720 GB/s

Weather and Climate on GPUs



How to program GPUs

Application

Libraries



Directives



Programming
Languages



How is it done?

```
!$acc enter data async create (this)

!fixed size arrays allocation
!$acc enter data async create (this%c_reff_so, this%c_
!$acc enter data async create (this%c_rhoquer, this%c_
```

```
!allocate pointers
!$acc enter data async create (this%r_eff_so, this%r_
!$acc enter data async create (this%r_ pack_scalar_f64_kernel = (
```

```
!$acc data present(this)

!$acc parallel async
!$acc loop gang vector collapse (3)
DO k= 1 , ke
  DO i= 1, nproma
    DO j= 1, nhabits_ice
      this%r_eff_so(i,k,j) = 0.0_w
      this%r_eff_th(i,k,j) = 0.0_w
      this%r_mean_th(i,k,j) = 0.0_w
    END DO
  END DO
END DO
!$acc end parallel
```

```
# reset fields
with computation(FORWARD), interval(...):
  zfix = 0
  sum0 = 0.0
  sum1 = 0.0
with computation(PARALLEL), interval(...):
  lower_fix = 0.0
  upper_fix = 0.0
# fix_top:
with computation(BACKWARD):
  with interval(1, 2):
    if q[0, 0, -1] < 0.0:
      q = (
        q + q[0, 0, -1] * dp[0, 0, -1] / d
      ) # move enough mass up so that the t
  with interval(0, 1):
    if q < 0:
      q = 0
      dm = q * dp
# fix_interior:
with computation(FORWARD), interval(1, -1):
  # if a higher layer borrowed from this one, account for that here
  if lower_fix[0, 0, -1] != 0.0:
    q = q - (lower_fix[0, 0, -1] / dp)
  if q < 0.0:
    zfix += 1
    if q[0, 0, -1] > 0.0:
      # Borrow from the layer above
      dm = q * dp
```

```
None
if cp is None
else cp.RawKernel(
  r"""
extern "C" __global__
void pack_scalar_f64(const double* i_sourceArray,
  const int* i_indexes,
  const int i_nIndex,
  const int i_offset,
  double* o_destinationBuffer)
{
  int tid = blockDim.x * blockIdx.x + threadIdx.x;
  if (tid>=i_nIndex)
  {
    return;
  }

  o_destinationBuffer[i_offset+tid] = i_sourceArray[i_indexes[tid]];
}

""",
"pack_scalar_f64",
)
)
```

How to program GPUs

Libraries



Directives



Programming
Languages



Lab Exercises

01-GPU-programming-cupy.ipynb

- Introduction to GPU programming using a high-level programming language

Remarks

- When running a GPU notebook, you may experience this error:

`cupy_backends.cuda.api.runtime.CUDARuntimeError: cudaErrorDevicesUnavailable: all CUDA-capable devices are busy or unavailable`

Don't worry, it's not your fault! Just restart the kernel and the error should disappear. If it persists, reach out to us.

- To let multiple tasks access the same GPU: `export CRAY_CUDA_MPS=1`