

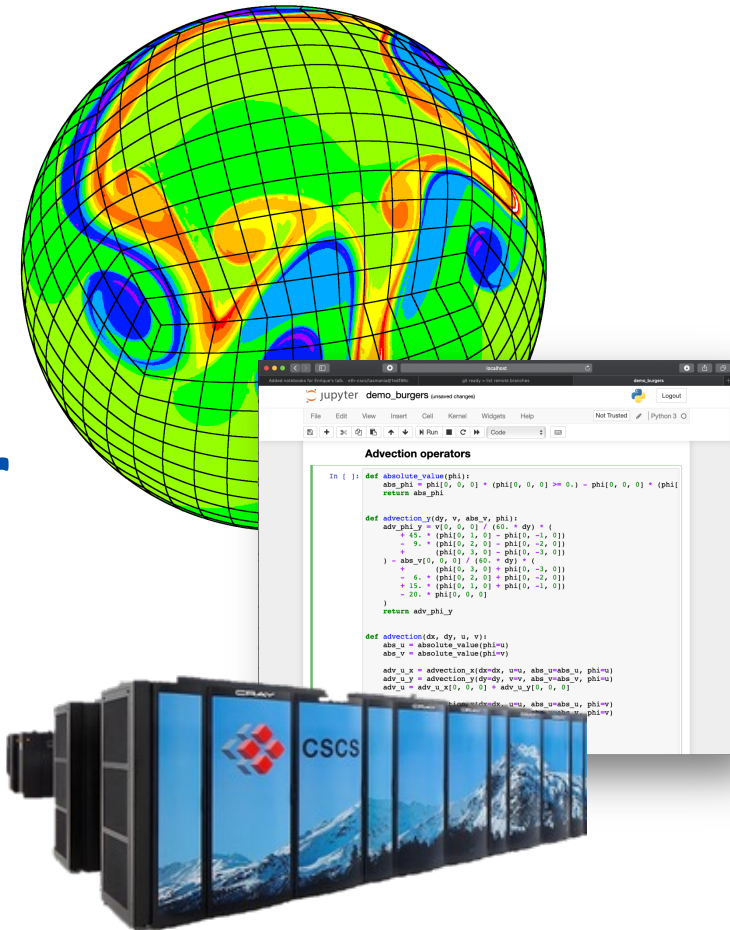
High Performance Computing for Weather and Climate (HPC4WC)

Content: High-Level Programming

Lecturer: Oliver Fuhrer

Block course 701-1270-00L

Summer 2024



Supercomputer Architecture

(Numbers are for Piz Daint and vary from system to system)

Day 1

- Single core performance
- Caches



Core
12/socket



Socket
2/node



Node
4/blade



Blade
48/cabinet

Day 3

- Multi-node performance
- Distributed memory parallelism
- MPI



Cabinet
40/system



System

I/O

Day 4

- Hybrid node architectures
- Graphics processing units (GPUs)
- CuPy

Future of HPC in Weather and Climate?

Yesterday

x86 CPU

MPI

Fortran

OpenMP

Today

GPU

C++

Python

mpi4py

CuPy

Tomorrow

Specialized hardware?

FPGA?

Domain-specific languages?

GT4Py?


ASIC?

ML?

Learning Goals

- Understand what a domain-specific language (DSL) is.
- Understand how a DSL helps in writing hardware-agnostic and maintainable code without sacrificing performance.
- Be able to apply a DSL to a stencil program from a weather and climate model.

Typical Workflow



Possible Scenarios

What if...

...we want to introduce a modification at the algorithmic/numerical level?

...our application has a broad user community and it must run efficiently on a variety of platforms?

...our code consists of thousands (if not millions) lines of code?

The explosion of hardware architectures made this development model obsolete!

A Real-Case Example: COSMO

- Limited-area model developed by the **C**onsortium for **S**mall-Scale **M**odeling.
- Run operationally by 8 national weather services and used by several academic institutions as a research tool.
- Three target architectures: x86 CPUs, NVIDIA GPUs and NEC vector CPUs
- Around 400K lines of F90 code and 90K lines of C/C++ code.
- Cost of porting the full code base to GPU: approx. 20-30 programmer-years!

Separation of Concerns

Domain expert

Answer scientific research questions

Declarative programming style:

Focus on **what** you want to do

Common data access interface:

e.g. `data[i, j, k]`

Computation kernels:

Calculations for a single grid point

Individual operators (“grains”)

Performance expert

Write optimized code for target platform

Imperative programming style:

Focus on **how** to do it

Storage and memory allocation:

e.g. C-layout vs F-layout

Control structure (e.g. for loops):

Optimized data traversal

Final computation:

Detect and exploit parallelism b/w grains

Overarching Goals (The 3 P's)

- **Productivity**

Easy to implement.

Easy to **read**.

Easy to **maintain**.

- **Performance**

Is **fast**.

- **Portability**

Single **hardware-agnostic** application code.

Runs efficiently on **different hardware** targets.

Domain Specific Languages (DSLs)

- Programming language tailored for a specific class of problems.
- Higher level of abstraction w.r.t. a general purpose language.
- Intended to be used by domain experts, who may not be fluent in programming.
- Abstractions and notations much aligned to concepts and rules from the domain.

Domain Specific Languages (DSLs)

- Programming language tailored for a specific class of problems.
- Higher level of abstraction w.r.t. a general purpose language.
- Intended to be used by domain experts, who may not be fluent in programming.
- Abstractions and notations much aligned to concepts and rules from the domain.
- Some examples:
 - Typesetting: LaTeX
 - Machine Learning: PyTorch, TensorFlow (Keras)
 - Scientific Computing: Kokkos, FEniCS, FreeFEM
 - Fluid Dynamics: OpenFOAM
 - Image Processing: Halide, Taichi
 - Stencils: Devito, Ebb, GT4Py

GT4Py

- High-performance implementation of a stencil kernel from a high-level definition.
- GT4Py is a domain specific **library** which exposes a domain specific **language** (GTScript) to express the stencil logic.
- GTScript is embedded in Python (**eDSL**).
 - Legal Python syntax and (almost) legal Python semantics.
- GT4Py = **GridTools For Python**
 - Harnessing the C++ GridTools ecosystem to generate native implementations of the stencils.
- Emphasis on tight integration with scientific Python stack.

What Does The GT4Py DSL Need?

```
import numpy as np
```

```
def laplacian_np(in_field):
```

```
    out_field = np.zeros_like(in_field)
```

```
    nx, ny, nz = in_field.shape
```

```
    for i in range(1, nx - 1):
```

```
        for j in range(1, ny - 1):
```

```
            for k in range(0, nz):
```

```
                out_field[i, j, k] = (
```

```
                    - 4 * in_field[i, j, k]
```

```
                    + in_field[i - 1, j, k]
```

```
                    + in_field[i + 1, j, k]
```

```
                    + in_field[i, j - 1, k]
```

```
                    + in_field[i, j + 1, k])
```

```
    return out_field
```

Input, output, and possibly temporary 3D fields

Nested loops iterating along both horizontal and vertical directions

Math operations

Indices and offsets

Definitions Function

```
import numpy as np
import gt4py as gt
from gt4py.gtscript import Field, PARALLEL, computation, interval
```

```
f64 = np.float64
```

```
def laplacian_defs(in_field: Field[f64], out_field: Field[f64]):
    with computation(PARALLEL), interval(...):
        out_field = (
            - 4. * in_field[ 0, 0, 0]
            +   in_field[-1, 0, 0]
            +   in_field[+1, 0, 0]
            +   in_field[ 0, -1, 0]
            +   in_field[ 0, +1, 0] )
```

Regular (named) function

Definitions Function

```
import numpy as np
import gt4py as gt
from gt4py.gtscript import Field, PARALLEL, computation, interval
```

```
f64 = np.float64
```

```
def laplacian_defs(in_field: Field[f64], out_field: Field[f64]):
    with computation(PARALLEL), interval(...):
        out_field = (
            - 4. * in_field[ 0, 0, 0]
            +   in_field[-1, 0, 0]
            +   in_field[+1, 0, 0]
            +   in_field[ 0, -1, 0]
            +   in_field[ 0, +1, 0] )
```

Input and output fields
(with type annotations)

Definitions Function

```
import numpy as np
import gt4py as gt
from gt4py.gtscript import Field, PARALLEL, computation, interval
```

```
f64 = np.float64
```

```
def laplacian_defs(in_field: Field[f64], out_field: Field[f64]):
```

```
    with computation(PARALLEL), interval(...):
```

```
        out_field = (
            - 4. * in_field[ 0, 0, 0]
            +   in_field[-1, 0, 0]
            +   in_field[+1, 0, 0]
            +   in_field[ 0, -1, 0]
            +   in_field[ 0, +1, 0] )
```

Any computation must be wrapped in a **with** construct which can be thought of as being a k-loop

Definitions Function

```
import numpy as np
import gt4py as gt
from gt4py.gtscript import Field, PARALLEL, computation, interval
```

```
f64 = np.float64
```

```
def laplacian_defs(in_field: Field[f64], out_field: Field[f64]):
    with computation(PARALLEL), interval(...):
        out_field = (
            - 4. * in_field[ 0, 0, 0]
            +   in_field[-1, 0, 0]
            +   in_field[+1, 0, 0]
            +   in_field[ 0, -1, 0]
            +   in_field[ 0, +1, 0] )
```

Iteration order in the
vertical direction :
PARALLEL, FORWARD,
BACKWARD

Definitions Function

```
import numpy as np
import gt4py as gt
from gt4py.gtscript import Field, PARALLEL, computation, interval
```

```
f64 = np.float64
```

```
def laplacian_defs(in_field: Field[f64], out_field: Field[f64]):
    with computation(PARALLEL), interval(...):
        out_field = (
            - 4. * in_field[ 0, 0, 0]
            +   in_field[-1, 0, 0]
            +   in_field[+1, 0, 0]
            +   in_field[ 0, -1, 0]
            +   in_field[ 0, +1, 0] )
```

Vertical region of application:
... = full column

Definitions Function

```
import numpy as np
import gt4py as gt
from gt4py.gtscript import Field, PARALLEL, computation, interval
```

```
f64 = np.float64
```

```
def laplacian_defs(in_field: Field[f64], out_field: Field[f64]):
    with computation(PARALLEL), interval(...):
```

```
        out_field = (
            - 4. * in_field[ 0, 0, 0]
            +   in_field[-1, 0, 0]
            +   in_field[+1, 0, 0]
            +   in_field[ 0, -1, 0]
            +   in_field[ 0, +1, 0] )
```

Each statement can be
thought of as being an ij-loop

Definitions Function

```
import numpy as np
import gt4py as gt
from gt4py.gtscript import Field, PARALLEL, computation, interval
```

```
f64 = np.float64
```

```
def laplacian_defs(in_field: Field[f64], out_field: Field[f64]):
```

```
    with computation(PARALLEL), interval(...):
```

```
        out_field = (
            - 4. * in_field[0, 0, 0]
            +   in_field[-1, 0, 0]
            +   in_field[+1, 0, 0]
            +   in_field[ 0, -1, 0]
            +   in_field[ 0, +1, 0] )
```

Neighboring points accessed
through offsets

Definitions Function

```
import numpy as np
import gt4py as gt
from gt4py.gtscript import Field, PARALLEL, computation, interval
```

```
f64 = np.float64
```

```
def laplacian_defs(in_field: Field[f64], out_field: Field[f64]):
```

```
    with computation(PARALLEL), interval(...):
```

```
        out_field = (
```

```
            - 4. * in_field[ 0, 0, 0]
```

```
            + in_field[-1, 0, 0]
```

```
            + in_field[+1, 0, 0]
```

```
            + in_field[ 0, -1, 0]
```

```
            + in_field[ 0, +1, 0] )
```

1st horizontal
dimension

2nd horizontal
dimension

Vertical
dimension

Neighboring points accessed
through offsets

Definitions Function

```
import numpy as np
import gt4py as gt
from gt4py.gtscript import Field, PARALLEL, computation, interval
```

```
f64 = np.float64
```

```
def laplacian_defs(in_field: Field[f64], out_field: Field[f64]):
    with computation(PARALLEL), interval(...):
        out_field = (
            - 4. * in_field[ 0, 0, 0]
            +   in_field[-1, 0, 0]
            +   in_field[+1, 0, 0]
            +   in_field[ 0, -1, 0]
            +   in_field[ 0, +1, 0] )
```

No for loops!
No return statement!

Compilation

- A stencil needs to be compiled for a given **backend**:

```
backend = "gt:cpu_ifirst"  
laplacian = gt.stencil(backend, laplacian_defs)
```

- Backends target different purposes, needs, and computer architectures:
 - Python: "numpy" (vectorized syntax);
 - C++: "gt:cpu_ifirst" (x86), "gt:cpu_kfirst" (MIC), "gt:gpu" and "cuda" (NVIDIA GPU).
- For non-Python backends, compilation consists of three steps:
 - 1) Generate optimized code for the target architecture (cached in .gt_cache).
 - 2) Compile the automatically generated code.
 - 3) Build Python bindings to that code.

Storages

- The compilation returns a callable object which can be invoked on GT4Py storages.
- Storages have optimal memory **strides**, **alignment** and **padding**.
- `gt.storage` provides functionalities to allocate storages ...

```
nx, ny, nz = 128, 128, 64
def_orig = (1, 1, 0)
out_field = gt.storage.zeros(
    backend, def_orig, (nx, ny, nz), dtype=f64 )
```

... and convert NumPy arrays into valid storages:

```
in_field = gt.storage.from_array(
    np.random.rand(nx, ny, nz), backend, def_orig, dtype=f64 )
```


Storages

- Storages can be accessed as NumPy arrays:

```
in_field[0, 0, 0] = 4.  
print(in_field[0, 0, 0])  
# Output: 4.0
```

Running

- Running computations is as simple as a function call:

```
laplacian(  
    in_field=in_field,  
    out_field=out_field,  
    origin=(1, 1, 0),  
    domain=(nx - 2, ny - 2, nz)  
)
```

Bindings b/w the symbols used
within the definitions fct.
and the arrays holding the data

Running


- Running computations is as simple as a function call:

```
laplacian(  
    in_field=in_field,  
    out_field=out_field,  
    origin=(1, 1, 0),  
    domain=(nx - 2, ny - 2, nz)  
)
```


Origin and extent of the
computation domain

- `out_field` now contains the results of the computation.

Region of application



Region of application



Weather and Climate on DSLs

- Several models (FV3, FVM and ICON) being ported to GT4Py
- Other approaches
 - COSMO (MeteoSwiss) dynamical was re-written in C++ using GridTools library.
 - E3SM (US DOE) using the Kokkos library for on-node parallelism.
 - LFric (UK MetOffice)
- Who knows what the future will bring...

Disadvantages of a DSL

- Lack of generality: A DSL is not a complete ontology!
- Debugging on the generated code.
- Cost of developing and maintaining the DSL compiler toolchain.

Conclusions

- High-level programming techniques hide the complexities of the underlying architecture to the end user.
- DSL allows to target multiple platforms without polluting the application code with hardware-specific boilerplate code.
- GT4Py is a Python framework to write performance portable applications in the weather and climate area. It ships with a DSL to write stencil computations.

Lab Exercises

01-GT4Py-motivation.ipynb

- Compare NumPy, CuPy and GT4Py on the sum-diff and Laplacian stencil (demo).

02-GT4Py-concepts.ipynb

- Digest the main concepts of GT4Py.
- Get familiar with writing, compiling and running stencils.
- Get insights on the internal data-layout of the storages.

04-GT4Py-stencil2d.ipynb

- Step-by-step porting of stencil2d.py to GT4Py.
- Write two alternative versions of stencil2d-gt4py-v0.py

References

Broad introduction to DSLs:

<https://www.jetbrains.com/mps/concepts/domain-specific-languages/>

GT4Py repository:

<https://github.com/GridTools/gt4py>

More in-depth introduction to GT4Py:

https://github.com/ai2cm/dsl_workshop