

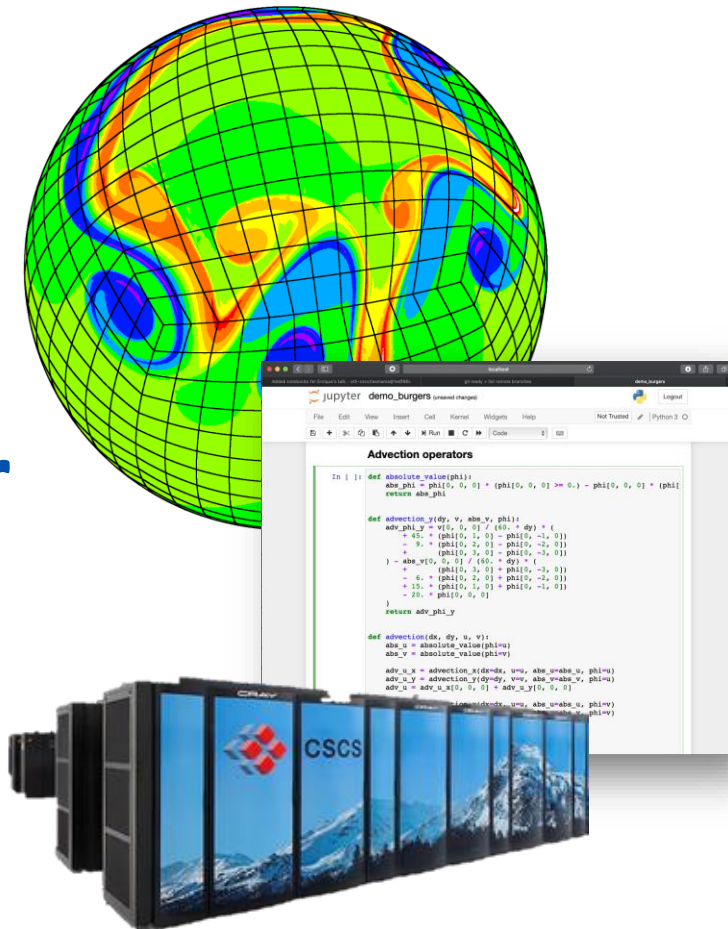
High Performance Computing for Weather and Climate (HPC4WC)

Content: Pace: A model in Python

Lecturer: Tobias Wicky

Block course 701-1270-00L

Summer 2022



Learning Goal

- See the DSL approach in action
- See some of the concepts we've learned applied in a real code

Who are we?



Oliver
Fuhrer



Johann
Dahm



Florian
Deconinck




Oliver
Elbert



Jeremy
McGibbon



Tobias
Wicky



Elynn Wu

Collaborators



CSCS



MeteoSwiss

ETH zürich

W
UNIVERSITY of
WASHINGTON

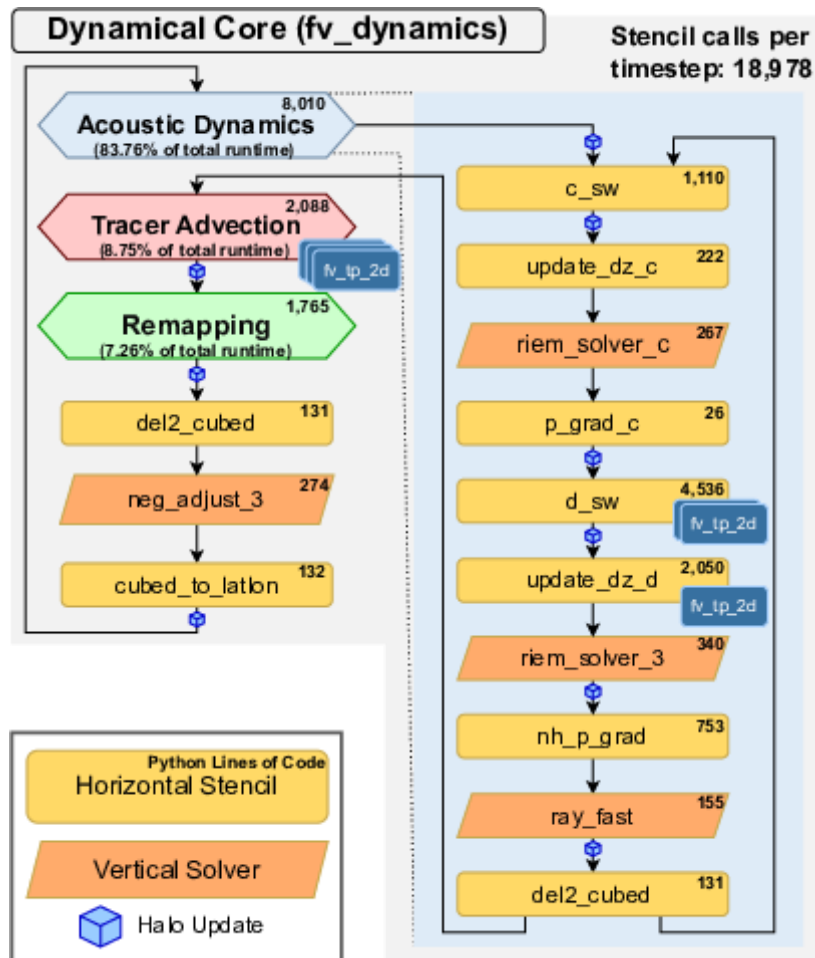
Pace: Python based FV3

GFDL Finite-Volume Cubed-Sphere Dynamical Core (FV3)




Finite volume transport on a cubed sphere
grid



- Integrated into several models, including
 - Operational weather models (Global Forecast System)
 - Next Generation Global Prediction System




Validation



Validation

Diff from init with heat source fix c48 6ranks: ucomp, z=40, t=6hr



Generated on 06/01/22 20:10:45


The Pace Model

Full program optimization


DSL coverage of all main numerical computation

Custom code for halo updates

New DSL concepts for FV3-specific motifs



The Pace Model



Why did you torture us for 4 days?

DSL still uses these concepts under the hood

```
def visit_Stencil(self, node: oir.Stencil, **kwargs: Any) -> oir.Stencil:
    write_before_read_tmps = {
        symbol
        for symbol, value in kwargs["syntable"].items()
        if isinstance(value, oir.Temporary)
    }
    horizontal_executions = node.iter_tree().if_isinstance(oir.HorizontalExecution)

    for horizontal_execution in horizontal_executions:
        accesses = AccessCollector.apply(horizontal_execution)
        offsets = accesses.offsets()
        ordered_accesses = accesses.ordered_accesses()

        def write_before_read(tmp: str) -> bool:
            if tmp not in offsets:
                return True
            if offsets[tmp] != {(0, 0, 0)}:
                return False
            return next(
                o.is_write and o.horizontal_mask is None
                for o in ordered_accesses
                if o.field == tmp
            )

        write_before_read_tmps = {
            tmp for tmp in write_before_read_tmps if write_before_read(tmp)
        }

    return super().visit_Stencil(node, tmps_to_replace=write_before_read_tmps, **kwargs)
```

Why did you torture us for 4 days?

There are things that we still need to do manually:

- Halo updates
- Memory size reduction

What does Python bring us?

```
class Physics:
```

```
...
```

```
    prepare_microphysics(physics_state)
```

```
    microph_state = physics_state.microphysics
```

```
    microphysics(microph_state)
```

GT4Py stencil-based

```
emulation_model = tf.keras.models.load_model("model.tf")
```

ML-based microphysics


```
emulation_dict = prepare_emulation_data(physics_state.microphysics)
```

```
predictions = emulation_model(emulation_dict)
```


```
model_outputs = unpack_predictions(predictions, emulation_model.output_names, ...)
```

ML-based Microphysics

Fortran liquid column sum ts=5



Python mph emulation liquid column sum ts=5



Why else is Python awesome?


The rich python ecosystem is valuable – new options for development

Testing is WAY easier!











You

What is the programming language you feel most comfortable in?

23 responses



Physical Parametrizations

	Microphysics	PBL & Turbulence	Sea-Ice	Shallow Convection	LSM	Radiation
Authors	 Mikael	 Chris Kung (NASA)	 Vera  Nadja  Nicolai	 Mikael  Chenwei  Langwen	 Safira	 Andrew
Scheme	GFDL Cloud Microphysics Scheme	GFS scale-aware EDMF PBL and Free Atmospheric Turbulence Scheme	GFS Sea Ice Scheme	GFS SAS-based Mass-Flux Scheme for Shallow convection	GFS Noah Land Surface Model	GFS RRTMG

Are we seeing what DSLs promise?

Overarching Goals (The 3 P's)

- **Productivity**
Easy to implement.
Easy to **read**.
Easy to **maintain**.
- **Performance**
Is **fast**.
- **Portability**
Single **hardware-agnostic** application code.
Runs efficiently on **different hardware** targets.


The Pace Model

Full program optimization

DSL coverage of all main
numerical computation

Custom code for halo
updates

New DSL concepts for FV3-
specific motifs



High-Level Programming 6

Physical Parametrizations

	Microphysics	PBL & Turbulence	Sea-Ice	Shallow Convection	LSM	Radiation
Authors	 Mikael	 Chris Kung (NASA)	 Wang, Roca, Minobe, Nicola	 Mikael, Chen, Lingren	 Safra	 Andrew
Scheme	GFDL Cloud Microphysics Scheme	GFS scale-aware EDMF PBL and Free Atmospheric Turbulence Scheme	GFS Sea Ice Scheme	GFS SAS-based Mass-Flux Scheme for Shallow convection	GFS Noah Land Surface Model	GFS RRTMG

High-Level Programming 13

```
stencil_config:  
  backend: numpy  
  rebuild: false  
  validate_args: true  
  format_source: false  
  device_sync: true  
initialization:  
  type: baroclinic  
performance_config:  
  performance_mode: false  
  experiment_name: c12_baroclinic  
comm_config:  
  type: read  
  config:  
    path: comm  
    rank: 0  
  nx_tile: 12
```